

Operating systems
Final examination
(part 1)
2009/2010

Student Name:

NIA:

1. (2pt) Suppose two processes enter the ready queue with the following properties: Process 1 has a total of 8 units of work to perform, but after every 2 units of work, it must perform 1 unit of I/O (so the minimum completion time of this process is 12 units). Assume that there is no work to be done following the last I/O operation. Process 2 has a total of 20 units of work to perform. This process arrives just behind P1. Show the resulting schedule for the shortest-job-first (preemptive) and the round-robin algorithms. Assume a time slice of 4 units for RR. What is the completion time of each process under each algorithm?

2. (2pt) Explain what will happen when the following C code is executed. Indicate how many times each of the letters "a", "b", "c" and "d" will be displayed and why. Explain your reasoning with the help of a tree diagram. *Also*, show two of the many distinct orders in which all of the letters might be displayed, depending on preemption/scheduling issues. Defend each with an explanation. (Note: assume that if a parent process finishes before a child process, this will not affect the child process, which is sometimes true and sometimes false depending on the operating system's setup.)

```
int i, status;

for (i = 0; i < 3; i++)
{
    if (fork() != 0)
    {
        printf("a");
        if (fork() == 0) {
            printf("c");
        } else {
            printf("d");
            exit(0);
        }
    } else {
        printf("b");
    }
}
```

3. (2pt) Consider the following C code that creates and joins with two threads. Assuming that the threads are scheduled completely before the parent process (i.e., have a higher priority), what will be the output from running this program? Be careful! There is a significant trick!

```
int a = 0;

void *print_fn(void *ptr)
{
    int tid = *(int *)ptr;
    int b = 0;

    a++; b++;
    printf("id: %d a: %d b: %d\n", tid, a, b);

    while (1); // Spin-wait here forever
}

int main()
{
    pthread_t t1, t2;
    int tid1 = 1;
    int tid2 = 2;
    int ret1, ret2;

    a++;
    printf("Parent says a: %d\n", a);
    ret1 = pthread_create(&t1, NULL, print_fn, (void *)&tid1);
    ret2 = pthread_create(&t2, NULL, print_fn, (void *)&tid2);

    if (ret1 || ret2) {
        fprintf(stderr, "ERROR: pthread_create failed\n");
        exit(1);
    }

    if (pthread_join(t1, NULL)) {
        perror("join of t1");
        exit(1);
    }
    if (pthread_join(t2, NULL)) {
        perror("join of t2");
        exit(1);
    }

    printf("Thread 1 and 2 complete\n");
}
```

4. (4pt) True/False

- a. Threads are cheaper to create than processes
- b. Kernel-scheduled threads are cheaper to create than user-level threads
- c. A blocking kernel-scheduled thread blocks all threads in the process
- d. Threads are cheaper to context switch than processes
- e. A blocking user-level thread blocks the process
- f. All kernel-scheduled threads of a process share the same virtual address space
- g. The operating system is not responsible for resource allocation between competing processes
- h. System calls do not change to privilege mode of the processor