
 <p>Universidad Carlos III de Madrid</p>	<p>Departamento de Informática Grado en Ingeniería Informática Sistemas Operativos</p> <p>Examen de la convocatoria ordinaria 20 de enero de 2011</p>	
-------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------

NOTES:

- Exam time: **3 hours**.
- You can NOT use either notes or calculators.
- The mobile phones have to be turned off during the exam (turned off and not only silenced)

NAME:

NIA:

GROUP:

Exercise 1. Theory and small questions [3 points]:

Q1. When does a process enter the zombie state?

- A.- When his father dies and he has not finished yet.
- B.- When his father dies without calling `wait` for him.
- C.- When the process dies and his father has not called `wait` for him.
- D.- When the process dies and his father has not finished yet.

Explain why.

Solution: C

A zombie is created only after a process dies and stays in that state until a parent picks up the exit code with the `wait` operation.


Q2. Which of the following scheduling policies is best suited for a time-sharing multiprogramming system?

- A.- Shortest Job First.
- B.- Round-Robin.
- C.- Priorities.
- D.- FIFO.

Explain why.

Solution: B. It is the only policy that guarantees interleaving of concurrently executing programs.

Q3. Is the following system in a deadlock? Why or why not?



Solution: No deadlock. P2 and P4 eventually release the resource R1 and R2 and P1 and P3 gain access to the resources.

Q4. Is the following system in a safe state? Why or why not?


	Allocation		Request		Availability	
	A	B	A	B	A	B
P1	0	0	5	1	1	5
P2	3	0	2	3		
P3	1	0	1	4		

Solution: safe state. P3 executes => Availability (2 5). P2 executes => Availability (5 5). Finally P1 executes.

Exercise 2 [2,5 points]:

1. Eratostene sieve is a algorithm of generating a series of consecutive prime numbers. In this exercise you are required to implement the Eratostene sieve with *processes and pipes*. The algorithm works in the following way :
 - A father creates N children.
 - The father connects to the child 0.
 - Child i connects to children i-1 and i+1 creating a pipeline.
 - The father generates a series of consecutive numbers 2, 3, 4, 5,..... and sends them one by one to child 0. The sequence is terminated by the number -1, which indicates that all the children have to finish.
 - Each child stores the first number to receive in a variable `local_prime`. Subsequently, if the number it receives is not a multiple of `local_prime`, it forwards it to the next child.
 - When a child receive a -1, it forwards it to the next children and terminates
 - The father waits for the termination of all children before it exits.

Answer:



```
#define N 10

int main()
{
    int p[N][2];
    int i;
    // CREATE PIPES
    for (i=0;i<N;i++)
        pipe(p[i]);

    for (i=0;i<N;i++){
        if (!fork()) {
            int local_prime, y;
            close(p[i][1]);
            close(p[i+1][0]);



            // FIRST NUMBER TO BE RECEIVED IS A PRIME NUMBER
            read(p[i][0], &local_prime, 4);
            printf(" CHILD %d received PRIME NUMBER %d\n", getpid(), local_prime);

            while (1) {
                read(p[i][0], &y, 4);

                // FORWARD A NUMBER ONLY IF IT IS NOT MULTIPLE OF THE LOCAL PRIME
                NUMBER
                if ((y%local_prime) &&(i < N-1))
                    write(p[i+1][1], &y, 4);
                // FINISH WHEN RECEIVING -1
                if (y == -1)
                    break;
            }
            exit(0);
        }
    }

    // SENDING A STREAM OF NUMBERS
    for (i=2;i<N*10;i++)
        write(p[0][1],&i,4);
    // SENDING -1 TO TERMINATE
    i=-1;
    write(p[0][1],&i,4);

    // WAITING FOR THE CHILDREN TO TERMINATE
    for (i=0;i<N;i++){
        int status,pid;
        pid=wait(&status);
        printf(" CHILD %d terminated. %n", pid);
    }
}
```

 <p>Universidad Carlos III de Madrid</p>	<p>Departamento de Informática Grado en Ingeniería Informática Sistemas Operativos</p> <p>Examen de la convocatoria ordinaria 20 de enero de 2011</p>	
-------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------

Exercise 3 [2 points]:

An array of 100 elements is processed for 200 iterations in the following way. In each iteration, for each index between 1 and 98 the average of the item and its two neighbors items is calculated and stored in the same item, i.e. $v[i] = (v[i-1] + v[i] + v[i+1]) / 3$. The neighbor of index 0 is considered to be 99 and viceversa.

The program is optimized by executing it in two threads: one half in **thread0** and another half in **thread1**. In each iteration the calculations are first done in an auxiliary array **vaux** local to each thread, so that the calculated values do not overwrite the current values. At the end of the iteration the result is written to the **v** array.

Summarizing each thread performs the following operations:

1. Calculate the average in the auxiliary array **vaux**.
2. When the calculation is finished, wait for the other thread to finish, before copying the **vaux** to **v**. In this way it is avoided that one thread overwrites necessary values.
3. Once a thread copies the data to the array **v** it must wait that the other thread also finishes copying the data to **v** before starting a new iteration.

Below you are given the main program structure. You are required to add the necessary variables and synchronization at marks I, II, III, IV and V.

```
#define SIZE 100
#define NUMITER 200
pthread_attr_t attr;
pthread_t idth[2];
float v[SIZE];
```

I // **ADD THE NECESSARY VARIABLES**

```
void *thread0(void *num) {
    int i,j;
    float vaux[SIZE/2];
```



II Thread 0//ADD THE NECESSARY VARIABLES

```
for (j=0; j<NUMITER; j++) {  
    vaux [0]= (v[SIZE-1]+v[0]+v[1])/3;  
    for (i=1; i<SIZE/2 ; i++)  
        vaux [i]= (v[i-1]+v[i]+v[i+1])/3;
```

III Thread 0//ADD THE NECESSARY SYNCHRONIZATION OPERATIONS

```
    }  
  
    pthread_exit(0);  
}  
  
void *thread1(void *num) {  
    int i,j;  
    float vaux[SIZE/2];
```

IV Thread 1// ADD THE NECESSARY VARIABLES

```
for (j=0; j<NUMITER; j++) {  
  
    for (i=SIZE/2; i<SIZE-1 ; i++)  
        vaux [i-SIZE/2]= (v[i-1]+v[i]+v[i+1])/3;  
    vaux [SIZE/2-1]= (v[SIZE-2]+v[SIZE-1]+v[0])/3;
```

V Thread 1 // ADD THE NECESSARY SYNCHRONIZATION OPERATIONS AND COPY

THE DATA TO THE ORIGINAL ARRAY



```
    }

    pthread_exit(0);
}

int main(){

    int i;
    initArray(); // initialize the array values
    pthread_mutex_init (&mtx1, NULL);
    pthread_mutex_init (&mtx2, NULL);
    pthread_attr_init(&attr);
    pthread_create(&idth[0],&attr,thread0,NULL);
    pthread_create(&idth[1],&attr,thread1,NULL);
    for (i=0; i<2; i++)
        pthread_join(idth[i],NULL);
    return(0);
}
```

Solution

```
I
pthread_mutex_t mtx1,mtx2;
pthread_cond_t varcond1,varcond2;
int confin1=0, confin2=0;

II
int init1=0, init2=0;

III
pthread_mutex_lock (&mtx1);
    confin1++;
    //The one waiting will reinit confin1 for the next iteration
    if (confin1==1) init1=1;
    //Wait for the other to finish
    while (confin1!=2)
        pthread_cond_wait(&varcond1, &mtx1);
    for (i=0; i<SIZE/2 ; i++)
        v[i]=vaux[i];
    pthread_cond_signal(&varcond1);
    if (init1){
        confin1=0;
```



```
    init1=0;
}
pthread_mutex_unlock (&mtx1);

pthread_mutex_lock (&mtx2);
contfin2++;
if (contfin2==1) init2=1;
//Wait for the other to copy the data
while (contfin2!=2)
    pthread_cond_wait(&varcond2, &mtx2);
pthread_cond_signal(&varcond2);
if (init2){
    contfin2=0;
    init2=0;
}
pthread_mutex_unlock (&mtx2);
```

IV

```
int init1=0, init2=0;
```

V

Same as thread 0 except:

```
for (i=SIZE/2; i<SIZE ; i++)
    v[i]=vaux[i-SIZE/2];
```



Exercise 4 [2,5 points]:

In a UNIX file system the block size have 1024 bytes, the block pointers are 2 bytes long, and the inodes contain:

- 16 direct block pointers.
- 2 simple indirect block pointers.
- 2 double indirect block pointers.

Answer the following questions:

- a. How many disk accesses are necessary for reading a file of size 1 Mbytes?
Assume that once a block is accessed from disk it is kept in cache.
- b. What is the theoretical maximum size of a file?
- c. What is the maximum possible size of the whole file system?
- d. Which are the advantages and disadvantages of increasing the block size? Hint:
Refer to the question b and c and to the efficiency of disk space utilization.

 <p>Universidad Carlos III de Madrid</p>	<p>Departamento de Informática Grado en Ingeniería Informática Sistemas Operativos</p> <p>Examen de la convocatoria ordinaria 20 de enero de 2011</p>	
-------------------------------------------------------------------------------------------------------------------------------	------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------

Solution:

- a. Necessary read data blocks:

$$10 \text{ MB}/1024 = (10 \cdot 2^{20})/2^{10} = 10 \cdot 2^{10} \text{ bloques de datos}$$

Simple indirection blocks direct:

$$(1024/2) \cdot 1024 = (2^{10}/2) \cdot 2^{10} = 2^{19} = 512 \text{ KB}$$

Double indirection blocks direct:

$$(1024/2) \cdot (1024/2) \cdot 1024 = (2^{10}/2) \cdot (2^{10}/2) \cdot 2^{10} = 2^{28} = 256 \text{ MB}$$

Total:

$10 \cdot 2^{10}$ data blocks + 1 access to inode block, 2 accesses to simple indirect blocks + 1 access to double indirect blocks

- b. $(16 \cdot 2^{10}) + 2 \cdot ((2^{10}/2) \cdot 2^{10}) + 2 \cdot ((2^{10}/2) \cdot (2^{10}/2) \cdot 2^{10}) =$
 $16\text{KB} + 2 \cdot 512\text{KB} + 2 \cdot 256\text{MB} = 16\text{KB} + 1\text{MB} + 512\text{MB}$

- c. You can direct in the FS maximum 2^{16} blocks. The maximum size: $2^{16} \cdot 1\text{KB} =$
64 MB

- d. Consequences:

- Increases the file size
- Increases the device size of the file system
- Improves the performance (less disk accesses to a file)
- Increases the probability of internal fragmentation.