

# OPERATING SYSTEMS: INTRODUCTION AND BASIC CONCEPTS

Operating System Services

# Goals

2

- To understand what an operating system service is.
- To understand mechanisms participating in a system call.
- To know POSIX interface features.
- To know main services offered by POSIX.

# Running the Operating System

3

- Once startup is finished, the operating system only runs in response to interrupts.
  
- Operating System activated in response to:
  - ▣ A service request from a process.
  - ▣ An interrupt (peripheral or clock).
  - ▣ Hardware exception.

# Service activation

4

- A direct invocation to an OS routine leads to security problems:
  - How to perform changes in secure mode of operation?
- Using a software interrupt, OS activation becomes safer.
  - Library routine:
    - Machine instructions to prepare OS call.
    - Trap instruction.
    - Instructions processing results from OS call.

# OS services: System calls.

5

- Interface between applications and OS.
  - ▣ Generally available as functions in assembler.
  - ▣ Currently also in high-level languages (C, C++, ...).
  
- Typical services from OS.
  - ▣ Process management.
  - ▣ Thread management.
  - ▣ Signals and timers management.
  - ▣ Memory management.
  - ▣ Files and directories management.
  
- Example calls:
  - ▣ read: Allows reading data from file.
  - ▣ fork: Allows creating a new process.


# Call invocation

6

- Each **API** (*Applications Programmer Interface*) corresponds to some OS **service**.
  - ▣ Function is a wrapper for the code invoking the OS service.
- Includes the execution of a **trap** instruction to transfer control to operating system by generating an interrupt.
- Operating system handles interrupt and returns control to user program.

# Call invocation

7



# Service selection

8

- A single **trap** instruction and multiple services.
  - ▣ Need some mechanism for parameter passing from user process to kernel.
  
- As a minimum, a specification of desired service needs to be passed.
  - ▣ Usually a **numeric descriptor**.



# Parameter passing

9

- Three generic methods to pass parameters for system calls:
  - In **registers**.
  - In a **table in memory**, which address is passed to OS through a register.
  - Placing parameters in **program stack** and allow OS to extract them.
  
- Each OS provides its own **system calls**:
  - POSIX in UNIX and LINUX.
  - Win32 in Windows NT.

# Handling routine

10

- Handling routine must:
  - ▣ Retrieve parameters sent by the user process.
  - ▣ Identify service to be executed.
  - ▣ Determine address of service routine (indexing in table of service routines).
  - ▣ Transfer control to service routine.

# Call invocation

11

```
int read(int fd, char * b, int lon) {  
    int r;  
    __asm__(`  
        lw $a0, fd  
        lw $a1, b  
        lw $a2, lon  
        li $v0, 14  
        syscall  
        sw $v0, r`  
    );  
    return r;  
}
```

The diagram illustrates the mapping of the `syscall` instruction to the `READ_SYSCALL` trap. A yellow arrow points from the text **READ\_SYSCALL** to the `li $v0, 14` instruction, indicating that the value 14 is the system call number for reading. A blue arrow points from the text **TRAP** to the `syscall` instruction, indicating that this instruction triggers a trap to the kernel.

# Programmer interface

12

- Interface offers view as extended machine that operating system user has.
- Each operating system may offer one or several interfaces:
  - Linux: POSIX
  - Windows: Win32, POSIX

# POSIX standard

13

- Standard interface for operating systems from IEEE.
- **Goal:** applications portability for different platforms and operating systems.
- **It is NOT** an implementation.
  - ▣ It only defines an interface.
- A family of standards
  - ▣ 1003.1 OS basic services.
  - ▣ 1003.1a Extensions to basic services.
  - ▣ 1003.1b Real time extensions.
  - ▣ 1003.1c Threading extensions.
  - ▣ 1003.2 Shell and utilities.
  - ▣ 1003.2b Additional utilities.

# POSIX characteristics

14

- Short function names in lower case:
  - ▣ **fork**
  - ▣ **read**
  - ▣ **close**
- Functions usually return 0 on success or -1 on error.
  - ▣ Variable **errno**.
- Resources managed by operating system referenced through **descriptors** (integer numbers).

# Example: Running a command

15

```
#include <sys/types.h>
#include <stdio.h>
int main(int argc, char** argv) {
    pid_t pid;
    pid = fork();
    switch (pid) {
        case -1: /* error */
            exit(-1);
        case 0: /* child process */
            if (execvp(argv[1], &argv[1])<0) { perror("error"); }
            break;
        default:
            printf("Parent process");
    }
    return 0;
}
```

prog cat f1

# Fork service

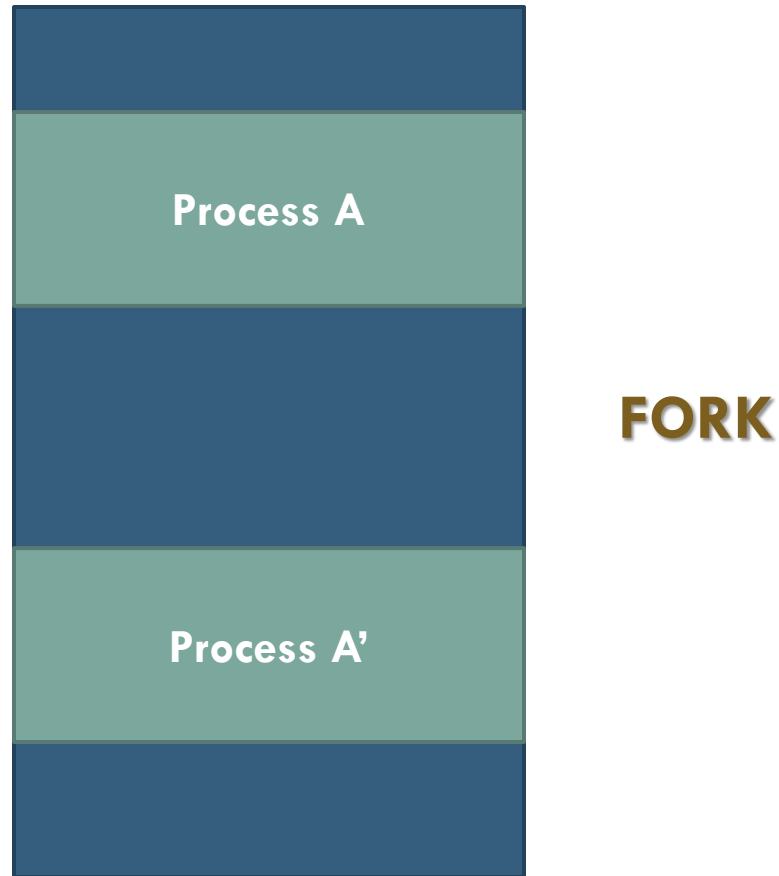
16

- `pid_t fork(void) ;`
- Duplicates process invoking the call.
- Parent process and child process go on running the same program.
- Child process inherits open files from parent process.
  - Open file descriptors are copied.
- Pending alarms are deactivated.
- Returns:
  - -1 on error.
  - In parent process: child process descriptor.
  - In child process: 0.



# Fork service

17



# Exec service

18

- Single service with multiple library functions.

```
int exec1(const char *path, const char *arg, ...);  
int execv(const char* path, char* const argv[]);  
int execve(const char* path, char* const argv[], char* const envp[]);  
int execvp(const char *file, char *const argv[])
```

- Changes current process image.

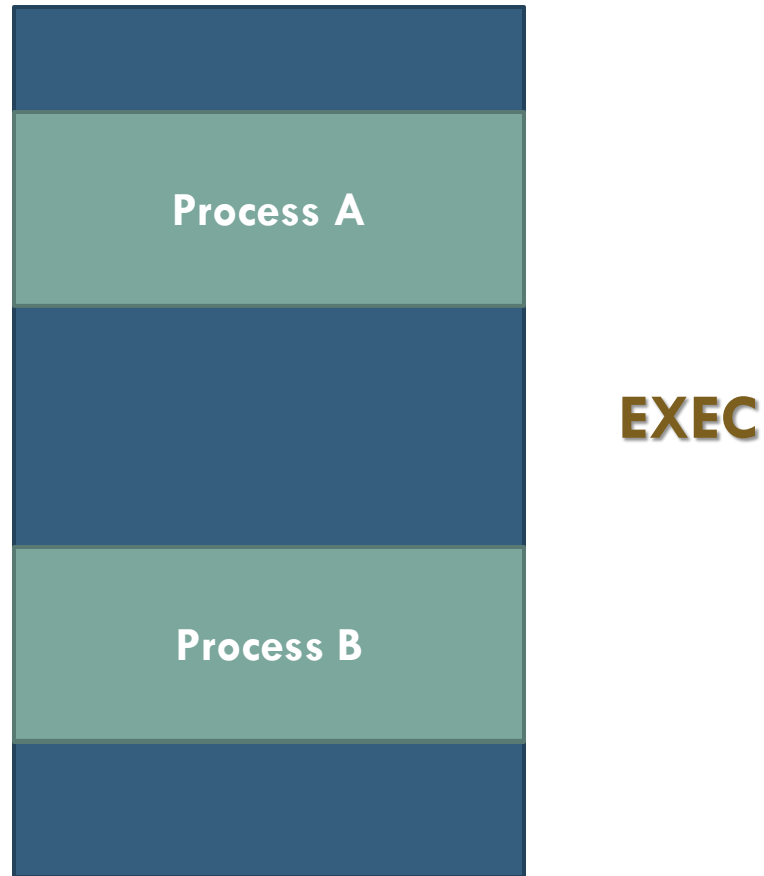
- **path**: path to executable file.
- **file**: Looks for the executable file in all directories specified by PATH.

- Description:

- Returns -1 on error, otherwise it does not return.
- The same process runs another program.
- Open files remain open.
- Signals with default action remain defaulted, signals with handler take default action.

# Exec service

19



# Exit service

20

- Finalizes process execution.

```
void exit(status) ;
```

- All open files descriptors are closed.
- All process resources are released.
- **PCB** (Process Control Block) is released.

# Example: Running a command

21

```
#include <sys/types.h>
#include <stdio.h>
int main() {
    pid_t pid;
    int status;

    pid = fork();
    if (pid == 0)    { /* child process */
        execlp("ls", "ls", "-l", NULL);
        exit(-1);
    }
    else            /* parent process */
        printf("Parent finalized\n");
    return 0; /* Invokes exit(0) */
}
```

# Generic operations on files

22

- **create:** Creates a file with name and attributes.
- **delete:** Deletes a file.
- **open:** Opens a file to allow access operations.
- **close:** Closes an open file.
- **read:** Reads data from open file to a memory buffer.
- **write:** Writes data to an open file from memory buffer.
- **position:** Moves pointer used to access file affecting subsequent operations.
- **control:** Allows manipulation of file attributes.

# POSIX services for files


23

- **Logical view:**
  - A file.
- Keeps pointer associated to every open file.
  - **Pointer:** position where next operation starts.
- Most operations work with file descriptors:
  - An integer number from 0 to 64K
  - Obtained upon opening file.
  - Rest of operations use file descriptors.
- Predefined descriptors:
  - 0: standard input.
  - 1: standard output.
  - 2: error output.

# POSIX services for files

24

- Every process has associated an open files table.
- When a process is duplicated (fork):
  - ▣ Duplicates open files table.
  - ▣ Shares intermediate table with i-nodes and positions.



- **Protection:**
  - ▣ owner group world
  - ▣ rwx rwx rwx
- **Examples:** 755 is `rwxr-xr-x`



# Files, directories and POSIX services

25

- File types:
  - ▣ Regular.
  - ▣ Directory.
  - ▣ Special.
- Names for files and directories:
  - ▣ Full name (starts with /)
    - `/usr/include/stdio.h`
  - ▣ Name relative to current directory (does not start with /)
    - `stdio.h` assuming `/usr/include` is current directory.
  - ▣ `.` and `..` entries can be used to form paths:
    - `../include/stdio.h`

# CREAT – Create file

26

## □ Servicio:

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
int creat(char *name, mode_t mode);
```

## □ Arguments:

- **name** File name
- **mode** Access rights bits.

## □ Returns:

- Return file descriptor or -1 upon error.

# CREAT – create file

27

- Description:
  - ▣ File is open for writing.
  - ▣ If not existing, creates an empty file.
    - `UID_owner = UID_effective`
    - `GID_owner = GID_effective`
  - ▣ If existing, truncates without changing access rights bits.

- Examples:

```
fd = creat("data.txt", 0751);
```

```
fd = open("data.txt",  
          O_WRONLY | O_CREAT | O_TRUNC, 0751);
```

# UNLINK – Erase file

28

## □ Service:

```
#include <unistd.h>
int unlink(const char* path);
```

## □ Arguments:

- **path** file name

## □ Returns:

- Returns 0 or -1 upon error.

## □ Description:

- Decrements link counter. If counter is 0, erases file and releases resources.

# OPEN – Open a file

29

## □ Service:

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
int open(char *name, int flag, ...);
```

## □ Arguments:

- **name** file name.

- **flags** options for opening:

- **O\_RDONLY** Read only.
- **O\_WRONLY** Write only.
- **O\_RDWR** Read/write
- **O\_APPEND** Access pointer moves to file end.
- **O\_CREAT** If existing has no effect. If not existing creates.
- **O\_TRUNC** Truncates if open for writing.

# Open – Opening a file

30

## □ Returns:

- ▣ A file descriptor or -1 upon error.

## □ Examples:

```
fd = open("/home/peter/data.txt");
```

```
fd = open("/home/peter/data.txt",  
          O_WRONLY | O_CREAT | O_TRUNC, 0750);
```

# CLOSE – Closing a file

31

- Service:

```
int close(int fd);
```

- Arguments:

- **fd** file descriptor.

- Returns:

- Zero or -1 upon error.

- Description:

- Process loses its link with the file.

# READ – Reading from a file

32

- **Service:**

```
#include <sys/types.h>
ssize_t read(int fd, void *buf, size_t n_bytes);
```

- **Arguments:**

- **fd** File descriptor.
- **buf** Buffer for data storage.
- **n\_bytes** Number of bytes to be read

- **Returns:**

- Number of bytes effectively read or -1 upo error.

- **Description:**

- Transfers **n\_bytes**.
- Can read less bytes when end of file is reached or interrupted by a signal.
- After reading the file pointer is incremented with the number of bytes effectively read.



# WRITE – Writing to a file

33

## □ Service:

```
#include <sys/types.h>
ssize_t write(int fd, void *buf, size_t n_bytes);
```

## □ Arguments:

- **fd** File descriptor.
- **buf** Buffer with data to be written.
- **n\_bytes** Number of bytes to be written.

## □ Returns:

- Number of bytes effectively written or -1 upon error.

## □ Description:

- Transfers **n\_bytes**.
- It may write less data than requested in file maximum size is reached or interrupted by a signal
- After writing file pointer is incremented with the number of bytes effectively written.
- If end of file is rebased, the file size is increased.

# LSEEK – Moving the file pointer

34

## □ Service:

```
#include <sys/types.h>
#include <unistd.h>
off_t lseek(int fd, off_t offset, int whence);
```

## □ Arguments:

- **fd**                      File descriptor.
- **offset**                Offset from base position.
- **whence**                Base position for offset.

## □ Returns:

- New position or -1 upon error.

## □ Description:

- Repositions pointer associated to a `fd`
- New position computation:
  - **SEEK\_SET** position = offset
  - **SEEK\_CUR** position = current position + offset
  - **SEEK\_END** position = file size + offset

# FNCTL – Attribute modification

35

- Service:

```
#include <sys/types.h>
int fnctl(int fildes, int cmd /* arg*/ ...);
```

- Arguments:

- **fildes**                      File descriptor
- **cmd**                        Command to modify attributes.

- Returns:

- 0 on success or -1 upon error.

- Description:

- Modifies attributes for an open file.

# DUP – Duplicate a file descriptor

36

## □ Service:

```
int dup(int fd);
```

## □ Arguments:

- `fd` file descriptor

## □ Returns:

- A file descriptor sharing all the properties of `fd` or -1 upon error.

## □ Description:

- Creates a new file descriptor having in common with the previous one:
  - Accesses to the same file.
  - Shares the same position pointer.
  - Access mode is identical.
- New descriptor gets the lowest available numeric value.

# FTRUNCATE – Space allocation for a file

37

## □ Service:

```
#include <unistd.h>
int ftruncate(int fd, off_t length);
```

## □ Arguments:

- **fd** File descriptor.
- **length** New file size.

## □ Returns:

- Return 0 or -1 upon error.

## □ Description:

- New file size is **length**.
- If **length** is 0 file is truncated.

# STAT – Information on a file

38

## □ Service:

```
#include <sys/types.h>
#include <sys/stat.h>
int stat(char *name, struct stat *buf);
int fstat(int fd, struct stat *buf);
```

## □ Arguments:

- **name**      File name.
- **fd**        File descriptor.
- **buf**        Pointer to object of type **struct stat**
- File information stored in buf.

## □ Returns:

- 0 on success or -1 upon error.

# STAT – Information on a file

39

## □ Description:

- ▣ Gets information about a file and stores in object of type `struct stat`:

```
struct stat {  
    mode_t    st_mode;    /* file mode*/  
    ino_t     st_ino;     /* i-node */  
    dev_t     st_dev;     /* device */  
    nlink_t   st_nlink;   /* number of links */  
    uid_t     st_uid;     /* owner UID */  
    gid_t     st_gid;     /* owner GID */  
    off_t     st_size;    /* number of bytes*/  
    time_t    st_atime;   /* last access */  
    time_t    st_mtime;   /* last modification */  
    time_t    st_ctime;   /* last data modification */  
};
```

# STAT – Information on a file

40

- Check file type in `st_mode`:

`S_ISDIR(s.st_mode)` Is directory?

`S_ISCHR(s.st_mode)` Is special character file?

`S_ISBLK(s.st_mode)` Is special block file?

`S_ISREG(s.st_mode)` Is regular file?

`S_ISFIFO(s.st_mode)` Is pipe or FIFO?



# UTIME – Altering date attributes

41

## □ Service:

```
#include <sys/stat.h>
#include <utime.h>

int utime(char *name, struct utimbuf *times);
```

## □ Arguments:

- **name**           File name.
- **times**           Structure with last access and modification dates.
  - **time\_t actime**           Access date.
  - **time\_t mctime**          Modification date.

## □ Returns:

- Returns zero or -1 upon error.

## □ Description:

- Change dates for last acces and last moficiation with values of structure **struct utimbuf**

# Example: Copying a file

42

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <stdio.h>

#define BUFSIZE 512

main(int argc, char **argv) {
    int fd_ent, fd_sal;
    char buffer[BUFSIZE];
    int n_read;

    /* open input file */
    fd_ent = open(argv[1],
                  O_RDONLY);
    if (fd_ent < 0) {
        perror("open");
        exit(-1);
    }

    /* create output file */
    fd_sal = creat(argv[2], 0644);
    if (fd_sal < 0) {
        close(fd_ent);
        perror("open");
        exit(-1);
    }
```

# Example: Copying a file

43

```
/* reading loop */
while ((n_read = read(fd_ent, buffer, BUFSIZE)) > 0) {
    /* write buffer on output file*/
    if (write(fd_sal, buffer, n_read) < n_read) {
        perror("write2");
        close(fd_ent); close(fd_sal);
        exit(-1);
    }
}

if (n_read < 0) {
    perror("read");
    close(fd_ent); close(fd_sal);
    exit(-1);
}

close(fd_ent); close(fd_sal);
exit(0);
}
```

# POSIX services for directories

44

- **Logical view:**
  - A directory is a file with records of structure DIR.
  - It can be operated as regular file for reading.
    - Do not write to it with regular writing calls.
- DIR structure:
  - **d\_ino;** // *i-node*
  - **d\_off;** // *Position in file of element in directory*
  - **d\_reclen;** // *Directory size.*
  - **d\_type;** // *Element type*
  - **d\_name[0];** // *File name **of variable length***
- **CARE NOTE:** Variable length records cannot be manipulated as fixed size records.
- **Solution:** System calls to improve directories.

# POSIX services for directories

45

- ❑ **DIR \*opendir(const char \*dirname);**
  - ❑ Open a directory and return a pointer of type **DIR** to the beginning.
- ❑ **int readdir\_r(DIR \*dirp, struct dirent \*entry, struct dirent \*\*result);**
  - ❑ Read next directory entry and returns result in a **struct dirent**.
- ❑ **long int telldir(DIR \*dirp);**
  - ❑ Get current position of pointer within directory file.
- ❑ **void seekdir(DIR \*dirp, long int loc);**
  - ❑ Advance from current position to position specified by loc. Never goes backward.
- ❑ **void rewinddir(DIR \*dirp);**
  - ❑ Reset file pointer and move it to the beginning.
- ❑ **int closedir(DIR \*dirp);**
  - ❑ Close directory file.

# Goal accomplished?

46

- ❑ To understand what an operating system service is.
- ❑ To understand mechanisms participating in a system call.
- ❑ To know POSIX interface features.
- ❑ To know main services offered by POSIX.

# OPERATING SYSTEMS: INTRODUCTION AND BASIC CONCEPTS

Operating System Services

# Projections in POSIX

48

```
void *mmap(void *direc, size_t len, int prot,  
int flags, int fd, off_t offset);
```

- Sets a projection from a process address space and a file.
  - ▣ Return memory address where file was projected.
  
- ▣ **direc**: address where projection is performed. If NULL OS selects one.
- ▣ **len**: specifies number of bytes to project.
- ▣ **prot**: Protection bits for the area.
- ▣ **flags**: Properties for the region.
- ▣ **fd**: File descriptor to be used in memory.
- ▣ **offset**: Initial offset on the file.




# Projections with mmap

49

- Protection types:
  - **PROT\_READ**: Can read.
  - **PROT\_WRITE**: Can write.
  - **PROT\_EXEC**: Can execute.
  - **PROT\_NONE**: Cannot access.
- Properties of a memory region:
  - **MAP\_SHARED**:
    - Shared region.
    - Modifications affect to file.
    - Child processes share region.
  - **MAP\_PRIVATE**:
    - Private region.
    - File is not modified.
    - Child processes get non-shared duplicates.
  - **MAP\_FIXED**:
    - File must be projected in an address specified by the call.

# POSIX Projection

50



# POSIX removing mapping

51

- `void munmap(void *direc, size_t len);`
  - Removes part of the process address space from address **direc** to address **direc + len**

# Example: Count number of blanks in file

52

```
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/mman.h>
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>

int main() {
    int fd;
    struct stat dstat;
    int i, n;
    char c,
    char * vec;

    fd = open("datos.txt", O_RDONLY);
    fstat(fd, &dstat);

    vec = mmap(NULL, dstat.st_size,
               PROT_READ, MAP_SHARED, fd, 0);
    close(fd);
    c = vec;
    for (i=0; i<dstat.st_size; i++) {
        if (*c == ' ') {
            n++;
        }
        c++;
    }
    munmap(vec, dstat.st_size);
    printf("n=%d, \n", n);
    return 0;
}
```

# Example: Copy a file

53

```
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/mman.h>
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>

int main() {
    int i, fd1, fd2;
    struct stat dstat;
    char * vec1, *vec2, *p, *q;

    fd1 = open("f1", O_RDONLY);
    fd2 = open("f2",
        O_CREAT|O_TRUNC|O_RDWR, 0640);
    fstat(fd1, &dstat);
    ftruncate(fd2, dstat.st_size);

    vec1= mmap (0, bstat.st_size,
        PROT_READ, MAP_SHARED, fd1, 0);
    vec2= mmap (0, bstat.st_size,
        PROT_READ, MAP_SHARED, fd2, 0);

    close(fd1); close(fd2);

    p=vec1; q=vec2;
    for (i=0; i<dstat.st_size; i++) {
        *q++ = *p++;
    }

    munmap (fd1, bstat.st_size);
    munmap (fd2, bstat.st_size);

    return 0;
}
```