

OPERATING SYSTEMS: PROCESSES

Threads and Processes

Content

2

- ❑ **Thread Concept.**
- ❑ Thread models.
- ❑ Design aspects.
- ❑ Threads in PThreads.
- ❑ Thread scheduling

Applications with concurrent processes

3

- A process includes a single thread of execution.
- Design of applications with multiple concurrent tasks:
 - A receiver process gets requests and launches a process per request.
 - A receiver process and a fixed set of request handling processes.

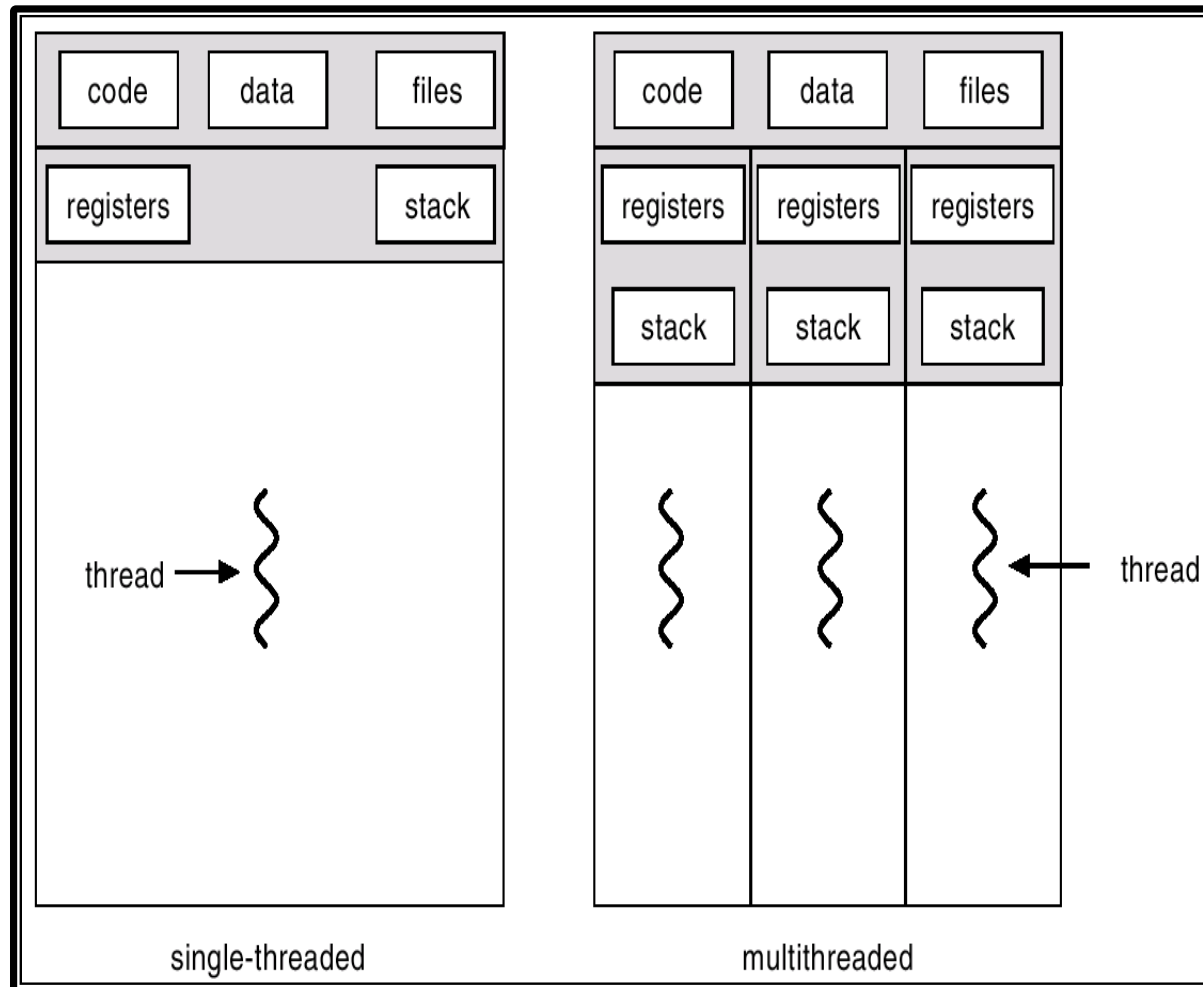
Performance of concurrent processes

4

- Time overhead in process creation and destruction.
- Time overhead due to context switching.
- Problems with sharing resources.

Threads

5



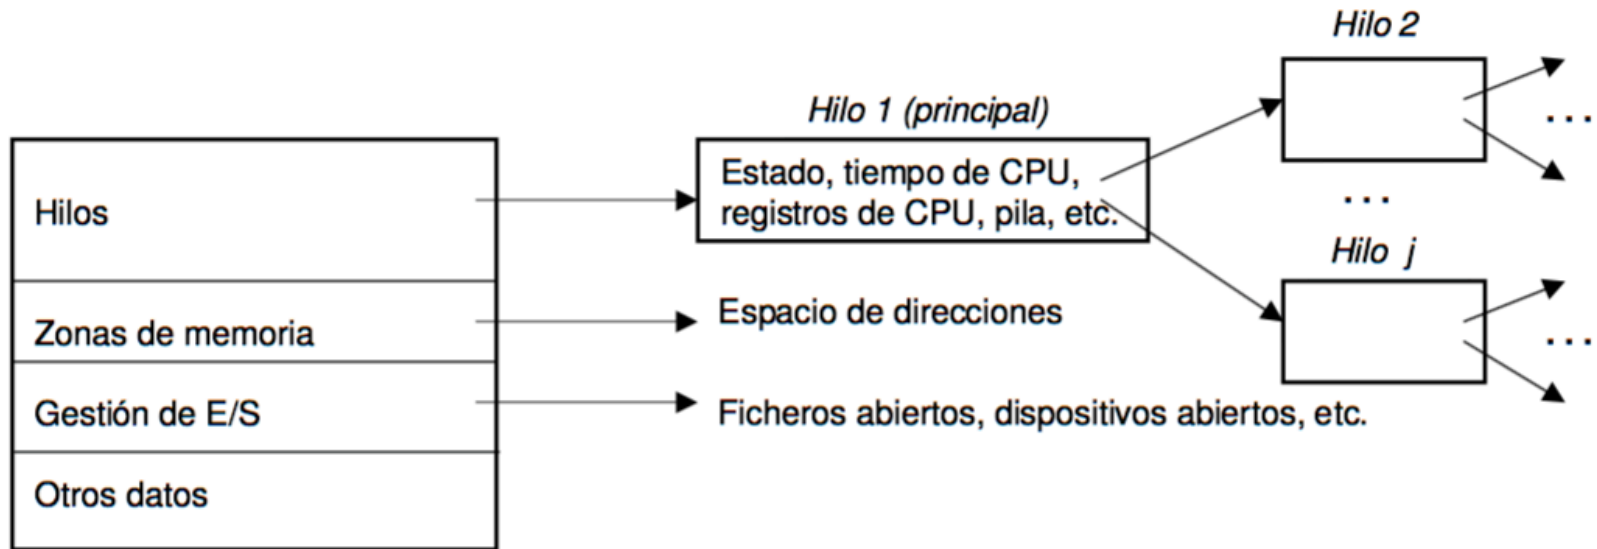
Threads

6

- Most modern OS provide processes with multiple instruction sequences or threads of control inside them.
- Used as the basic unit of CPU usage.
- Each one consists of:
 - ▣ Thread Id.
 - ▣ Program counter.
 - ▣ Registers set.
 - ▣ Stack.
- The share with the rest of threads in the process:
 - ▣ Memory map (code region, data region, shmem).
 - ▣ Open files.
 - ▣ Signals, semaphores, and timers.

BCP with threads

7



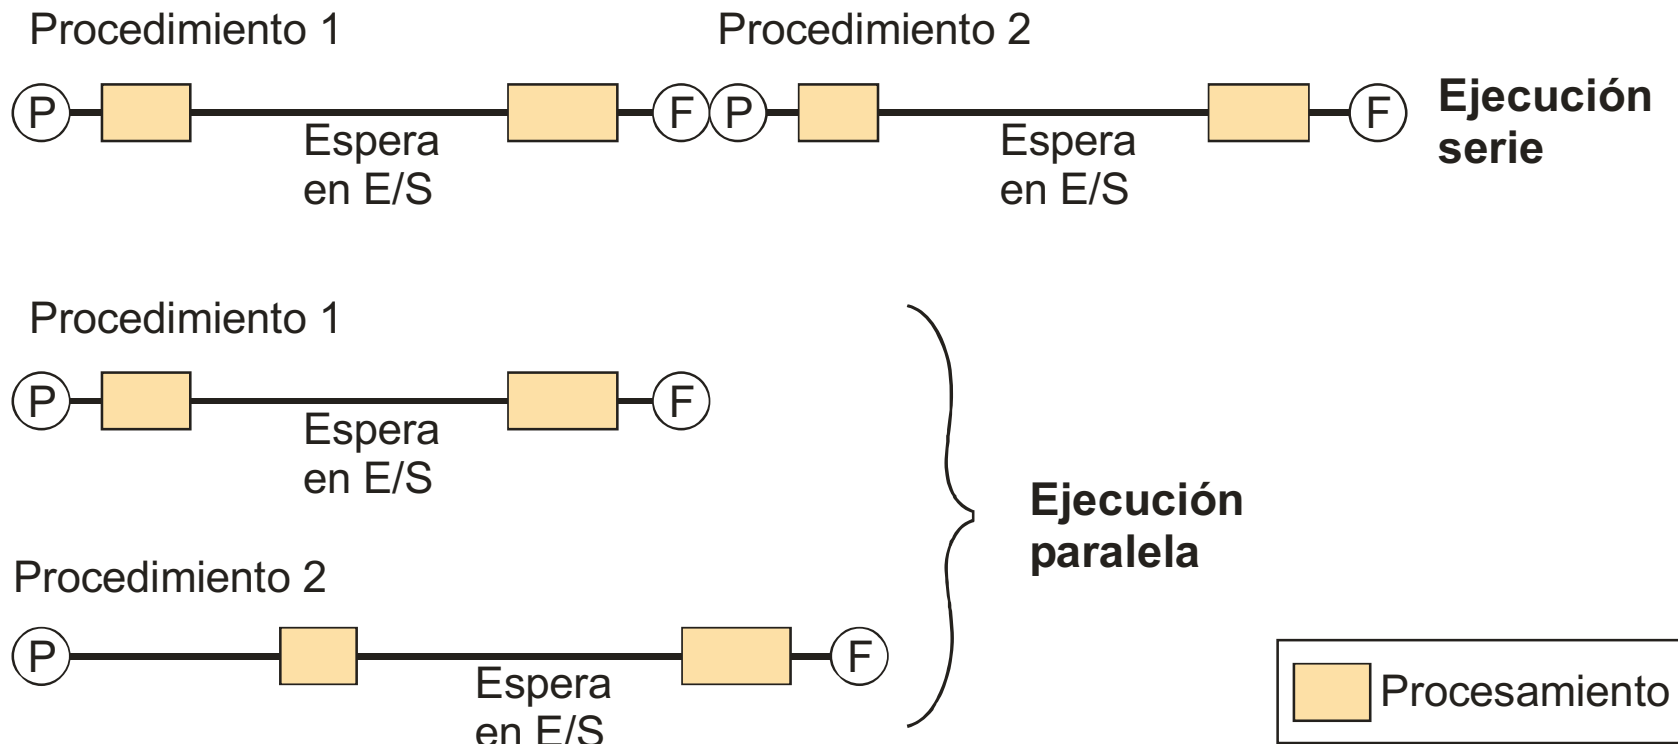
Benefits

8

- ❑ Response capacity.
 - ▣ Higher interactivity as user interactions are separated from processing tasks in different threads.
- ❑ Resource sharing.
 - ▣ Threads share most of the resources automatically.
- ❑ Resource economy.
 - ▣ Creating a process has a much higher overhead than creating a thread (e.g.: in Solaris the ratio is 30 to 1).
- ❑ Use in multiprocessor architectures.
 - ▣ Higher concurrency degree allocating different threads to different processors.
 - ▣ Most modern operating systems use thread as unit of scheduling.

Threads allow parellization of applications

9



Thread support

10

User space

ULT – User Level Threads

- ❑ Implemented in form of functions library in user space.
- ❑ Kernel has no knowledge about them.
 - ▣ No support in any form.
- ❑ Much faster, but some problems arise.
 - ▣ Blocking system calls.

Kernel space

KLT – Kernel Level Threads

- ❑ Kernel in charge of creating, scheduling and destroying them.
- ❑ Slightly slower as kernel needs to participate implying execution mode switching.
- ❑ In blocking system calls, only the involved thread is blocked.
- ❑ In SMP, multiple threads can run at a time.
- ❑ No support thread code in applications.
- ❑ Kernel may also use threads to perform its own tasks.

Content

11

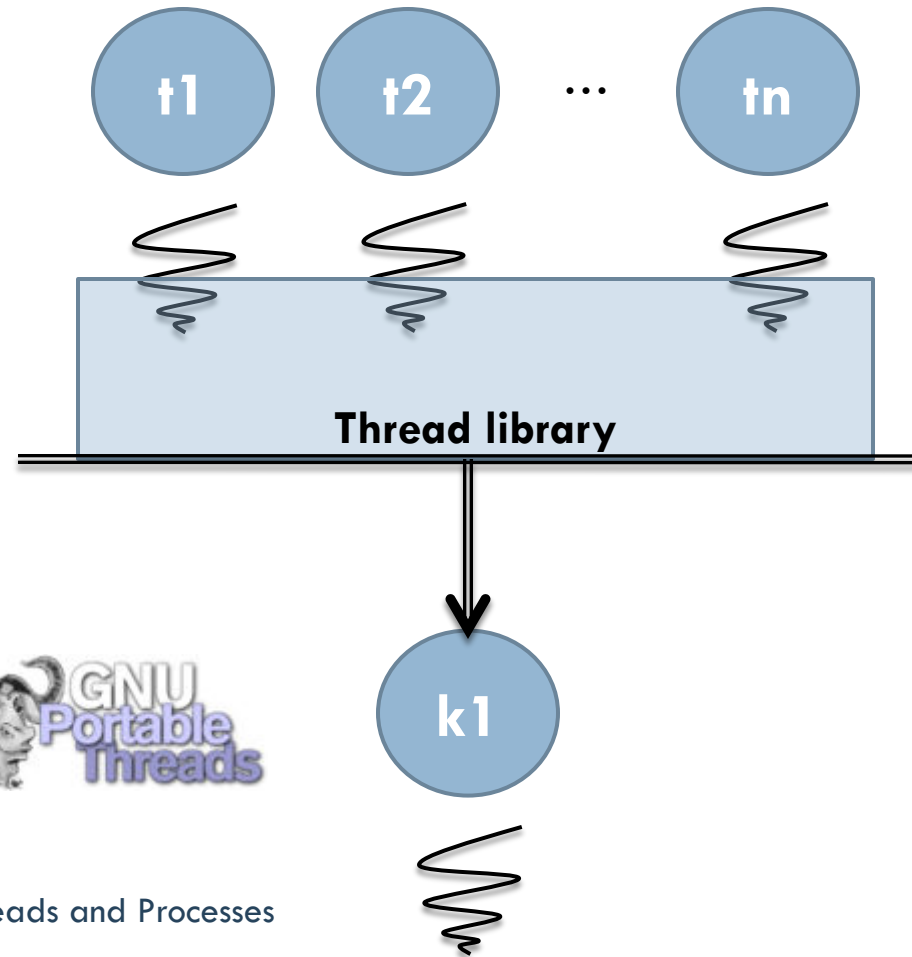
- Thread Concept.
- **Thread models.**
- Design aspects.
- Threads in PThreads.
- Thread scheduling

Multiple threads model:

Many to one

12

- ❑ Maps multiple user threads to a single kernel thread.
- ❑ Thread library in user space.
- ❑ Blocking call:
 - ❑ All threads are blocked.
- ❑ In multiprocessors, multiple threads cannot run at same time.




Multiple threads model:

Many to many

13

- This model multiplexes user threads into a fixed number of kernel threads.
- Operating System threads gets much more complex.
- Examples:
 - ▣ Solaris (before version 9).
 - ▣ HP-UX.
 - ▣ IRIX.




Multiple threads model:

One to one

14

- Maps a kernel thread to every user thread.
- Most implementations restrict the number of threads that can be created.
- Examples:
 - ▣ Linux 2.6.
 - ▣ Windows.
 - ▣ Solaris 9.



Content

15

- ❑ Thread Concept.
- ❑ Thread models.
- ❑ **Design aspects.**
- ❑ Threads in PThreads.
- ❑ Thread scheduling

Fork and exec calls

16

- In UNIX type systems. What to do if fork is called from a thread?
 - ▣ Duplicate process with all the threads.
 - Appropriate if exec is not going to be called immediately after to change process image.
 - ▣ Duplicate process with only the calling thread.
 - More efficient if exec is going to be called and all threads would be cancelled in any case.
- **Linux solution: Two versions of fork.**

Threads cancellation

17

- Situation when a thread notifies to others that the must terminate.
- Options:
 - Asynchronous cancellation: Forces immediate termination of thread.
 - Problems with resources allocated to thread.
 - Deferred cancellation: Thread checks periodically whether it should terminate.
 - Better approach.

Threads and request processing

18

- Applications receiving requests and processing them may make use of threads for handling.

- But:
 - ▣ Thread creation/destruction time is a delay (although lower than process creation/destruction).
 - ▣ A limit in the number of concurrent threads is not established.
 - ▣ If a request avalanche comes resources may be exhausted.

Thread Pools


19

- A *Thread Pool* is created and threads wait until requests arrive.

- Advantages:
 - ▣ Delay minimization: Thread already exists.
 - ▣ There is a established limit over the number of concurrent threads.

Software architectures based on threads

20



Content

21

- ❑ Thread Concept.
- ❑ Thread models.
- ❑ Design aspects.
- ❑ **Threads in PThreads.**
- ❑ Thread scheduling

Thread creation

22

- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*func)(void*), void *arg)`
 - Creates a thread and starts its execution.
 - **thread**: Must pass the address of a variable of type **pthread_t** used as handle.
 - **attr**: Must pass the address of a structures with attributes. NULL may be passed to specify default attributes.
 - **func**: Function with thread execution code.
 - **arg**: Pointer to thread parameter. Only one parameter can be passed.
- `pthread_t pthread_self(void)`
 - Returns thread identifier for the calling thread.

Waiting and termination

23

- `int pthread_join(pthread_t thread, void **value)`
 - ▣ Invoking thread waits until the thread identified by the handle has terminated.
 - ▣ **thread**: Handle to thread that must be waited to terminate.
 - ▣ **value**: Thread termination value

- `int pthread_exit(void *value)`
 - ▣ Allows a thread to terminate its execution, stating its termination status.
 - ▣ Termination status cannot be a pointer to a local variable.

Example: sum with threads

24

```
#include <stdio.h>
#include <pthread.h>

struct addparam{
    int n, m, r;
};
typedef struct addparam
    addparam_t;

void add(addparam_t * par) {
    int i;
    int sum=0;
    for (i=par->n;i<=par->m;i++)
    {
        sum+=i;
    }
    par->r=sum;
}

int main() {
    pthread_t th1, th2;
    sumapar_t s1 = {1,50,0};
    sumapar_t s2 = {51,100,0};

    pthread_create(&th1, NULL,
        (void*)_add, (void*)&s1);
    pthread_create(&th2, NULL,
        (void*)_add, (void*)&s2);

    pthread_join(th1, NULL);
    pthread_join(th2, NULL);

    printf("Total sum=%d\n",
        s1.r+s2.r);
}
```


Thread attribute

25

- Each thread has a set of attributes associated to it.
- Attributes represented by a variable of **pthread_attr_t** type.
- Attributes control:
 - ▣ Whether a thread is detached or joinable.
 - ▣ Size of the thread private stack.
 - ▣ Location of the thread stack.
 - ▣ Thread scheduling policy.

Attributes

26

- `int pthread_attr_init(pthread_attr_t * attr);`
 - ▣ Initializes a thread attribute structure.
- `int pthread_attr_destroy(pthread_attr_t * attr);`
 - ▣ Destroys a thread attribute structure.
- `int pthread_attr_setstacksize(pthread_attr_t * attr, int stacksize);`
 - ▣ Defines the stack size for a thread.
- `int pthread_attr_getstacksize(pthread_attr_t * attr, int *stacksize);`
 - ▣ Allows to obtain the size of the thread stack.

Hilos dependientes e hilos independientes

27

- `int pthread_attr_setdetachstate(pthread_attr_t *attr, int detachstate)`
 - ▣ Sets the termination state for a thread.
 - ▣ If **`detachstate == PTHREAD_CREATE_DETACHED`**
 - Thread releases its resources upon thread termination.
 - ▣ if **`detachstate = PTHREAD_CREATE_JOINABLE`**
 - Resources are not released automatically.
 - A call to **`pthread_join()`** is needed.
- `int pthread_attr_getdetachstate(pthread_attr_t *attr, int *detachstate)`
 - ▣ Allows to get the termination status.

Example: Detached threads

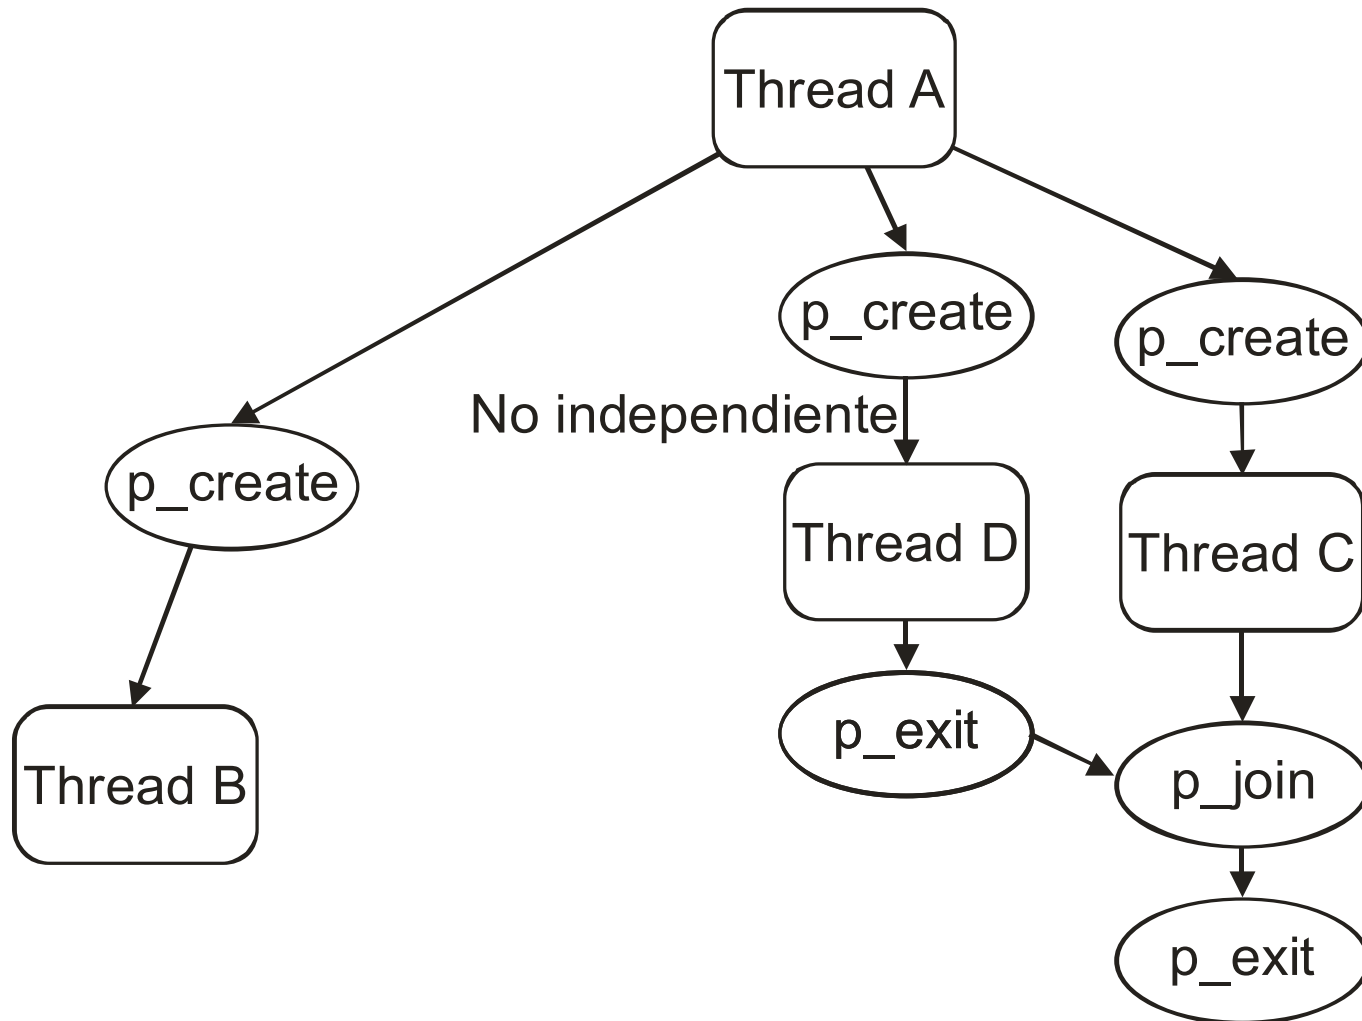
28

```
#include <stdio.h>
#include <pthread.h>
#define MAX_THREADS 10
void func(void) {
    printf("Thread %d \n", pthread_self());
    pthread_exit(0);
}

int main() {
    int j;
    pthread_attr_t attr;
    pthread_t thid[MAX_THREADS];
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);
    for(j = 0; j < MAX_THREADS; j++)
        pthread_create(&thid[j], &attr, func, NULL);
    sleep(5);
}
```

Example of threads hierarchy

29



Content


30

- ❑ Thread Concept.
- ❑ Thread models.
- ❑ Design aspects.
- ❑ Threads in PThreads.
- ❑ **Thread scheduling**

State of a process with threads

31

- Combination of the states of its threads:
 - ▣ If there is a thread executing -> Process executing
 - ▣ If no thread executing, but ready -> Process ready
 - ▣ If no threads blocked -> Process blocked



Thread scheduling into a process

- Based on priority model.
 - ▣ Segments of time not used.
- A thread will keep on running on the CPU until it goes to stopped or blocked state.
 - ▣ If you want to alternate threads in a process, it must be ensured and forced
 - ▣ For example using *sleep()*

Thread scheduling

33

- API to specify scheduling policy (PCS or SCS) on thread creation
 - ▣ PTHREAD_SCOPE_PROCESS: PCS scheduling
 - ▣ PTHREAD_SCOPE_SYSTEM: SCS scheduling
- SO can limit policy and numbers:
 - ▣ LINUX and MacOS only allows PCS for users (no kernel threads)

Thread scheduling API

34

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[]) {
    int i, scope;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* first inquire on the current scope */
    if (pthread_attr_getscope(&attr, &scope) != 0)
        fprintf(stderr, "Unable to get scheduling scope\n");
    else {
        if (scope == PTHREAD_SCOPE_PROCESS)
            printf("PTHREAD_SCOPE_PROCESS");
        else if (scope == PTHREAD_SCOPE_SYSTEM)
            printf("PTHREAD_SCOPE_SYSTEM");
        else
            fprintf(stderr, "Illegal scope value.\n");
    }
}
```

Thread scheduling API

35

```
/* set the scheduling algorithm to PCS or SCS */
pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
/* create the threads */
for (i = 0; i < NUM_THREADS; i++)
    pthread_create(&tid[i], &attr, runner, NULL);
/* now join on each thread */
for (i = 0; i < NUM_THREADS; i++)
    pthread_join(tid[i], NULL);
}

/* Each thread will begin control in this function */
void *runner(void *param)
{
    /* do some work ... */
    pthread_exit(0);
}
```

Summary

36

- A process may have several threads of execution.
- A multi-threaded application consumes less resources than a multi-process application.
- Each system has a thread support mode:
 - ▣ ULT versus KLT.
- PThreads is a user library for threading.
- Win32 offers kernel threads with support for Thread Pools.