

OPERATING SYSTEMS: PROCESS COMMUNICATION AND SYNCHRONIZATION

Concurrent Processes and problems in communication and synchronization

Contents

2

- **Concurrency.**
- Race conditions.
- Mutual exclusion and critical section.
- Semaphores.
- Producer-Consumer problem.
- Readers-Writers problem.

Concurrent process

3

- Two processes are **concurrent** when they are executed in a way that their **execution intervals overlap**.



There is concurrency



There is NO concurrency

Kinds of concurrency

4

- **Apparent concurrency:** More processes than processors
 - Processes are multiplexed in time.
 - Pseudoparallelism.

1 CPU



2 CPUs




Kinds of concurrency

5

- **Real concurrency:** Each process runs in a processor.
 - A **parallel execution** happens.
 - Also known as **real parallelism**..

4 CPUs



Concurrent programming models

6

- **Multiprogramming** with single processor.
 - ▣ OS in charge of dividing time among processes (appropriative/non-appropriative scheduling).
- **Multiprocessor.**
 - ▣ Combines real parallelism and pseudoparallelism.
 - Usually more processes than CPUs.
- **Distributed system.**
 - ▣ Several computers connected through a network.

Advantages of concurrent execution

7

- **Eases programming.**
 - Several tasks can be structured in separate processes.
 - **Web server:** A process in charge of handling each request.
- **Accelerates computation execution.**
 - Division of computations in processes run in parallel.
 - **Examples:** Simulations, Image processing, energy market, financial computations.
- **Improves application interactivity.**
 - Ability to separate processing tasks from user attention tasks.
 - **Example:** Print and edit.
- **Improve CPU utilization.**
 - I/O phases from one application are used for processing from other applications.

Kinds of concurrent processes

8

- Independent.
 - Processes run concurrently but without interaction.
 - No need for communication.
 - No need for synchronization.
 - Example: Two shells for two users run in two different terminals.
- Cooperating.
 - Processes run concurrently with some interaction.
 - May communication each other.
 - May be synchronized.
 - Example: Transaction server organized in receiver process and request handling processes.

Interactions among processes

9

- ❑ **Access to shared resources:**
 - ❑ Processes sharing a resource.
 - ❑ Processes competing for a resource.
 - ❑ **Example:** Request server with different processes writing to a log file.

- ❑ **Communication:**
 - ❑ Processes exchanging information.
 - ❑ **Example:** Requests receiver must pass information to a request handling process.

- ❑ **Synchronization:**
 - ❑ A process must wait for an event in another process.
 - ❑ **Example:** A user interface process must wait until all computation process have finished.


Contents

10

- Concurrency.
- **Race conditions.**
- Mutual exclusion and critical section.
- Semaphores.
- Producer-Consumer problem.
- Readers-Writers problem.

Race condition

11



Race condition

12

```
#include <stdio.h>
#include <pthread.h>

#define NUMTH 10
int total = 0;

void add() {
    int i,n;
    int sum = total;
    sum = sum + 100;
    n=rand()%5;
    for (i=0;i<n;i++)
        {printf(".");}
    total = sum;
}
```

```
int main() {
    pthread_t th[NUMTH];
    int i;
    for (i=0;i<NUMTH;i++) {
        pthread_create(&th[i],
            NULL,(void*)add, NULL);
    }

    for (i=0;i<NUMTH;i++) {
        pthread_join(th[i], NULL);
    }

    printf("Result=%d\n",
        total);
}
```

Result?

Result

13

```
[jdaniel@tucan ~]$ ./test2
.....Suma=200
[jdaniel@tucan ~]$ ./test2
.....Suma=600
[jdaniel@tucan ~]$ ./test2
.....Suma=500
[jdaniel@tucan ~]$ ./test2
.....Suma=300
[jdaniel@tucan ~]$ ./test2
.....Suma=600
[jdaniel@tucan ~]$ ./test2
.....Suma=600
[jdaniel@tucan ~]$ ./test2
.....Suma=500
[jdaniel@tucan ~]$ ./test2
.....Suma=600
[jdaniel@tucan ~]$ ./test2
.....Suma=600
[jdaniel@tucan ~]$ ./test2
.....Suma=600
[jdaniel@tucan ~]$ ./test2
.....Suma=500
```

- Each run gives a different result.
- Never gets the correct result.
- What's happening?

Possible sequences

14

`total = 0`

`sum = total`
`sum = sum + 100`

`total = sum`

`sum = total`
`sum = sum + 100`

`total = sum`

200

`total = 0`

`sum = total`
`sum = sum + 100`

`sum = total`
`sum = sum + 100`


`total = sum`

`total = sum`

100

Let's try again

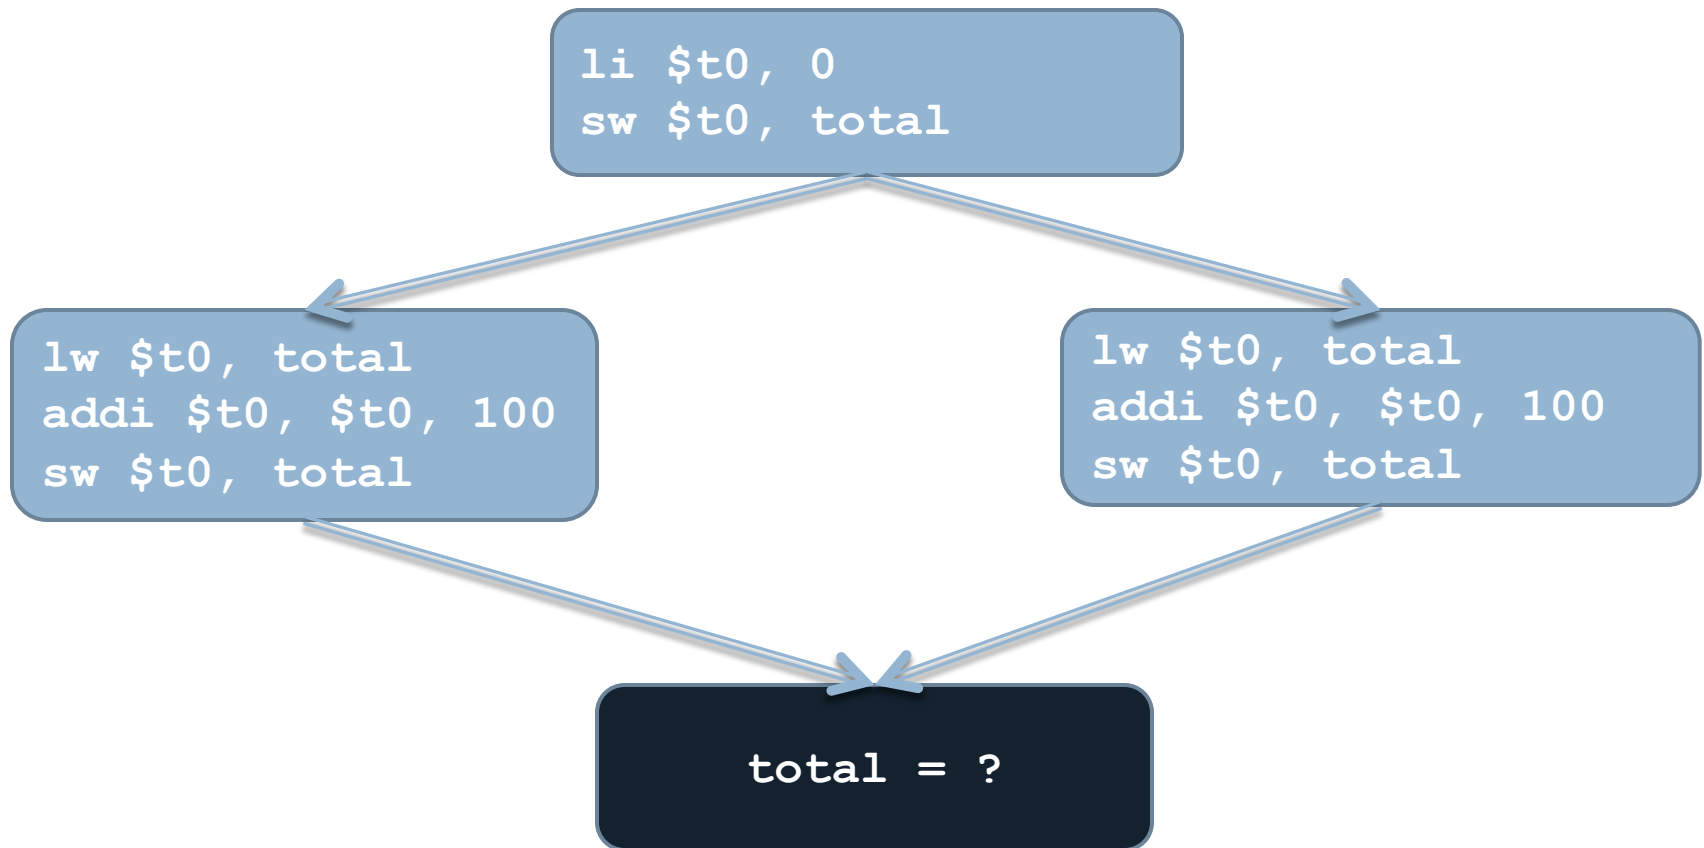
15



Is a race condition possible?

Machine instructions

16



Can this happen in a multiprocessor?

Race conditions

17

- Functioning of a process and its results must be independent from its relative execution speed with respect to other processes.
 - ▣ Necessary to guarantee that order of execution does not affect result.
- **Solution:** Achieve that a set of instructions can be executed atomically.

Mutual exclusion

Contents


18

- Concurrency.
- Race conditions.
- **Mutual exclusion and critical section.**
- Semaphores.
- Producer-Consumer problem.
- Readers-Writers problem.

Mutual exclusion

19

- **Critical section:** Code segment manipulating a resource that must be **atomically executed**.
- A resource is associated to a **management mechanism for mutual exclusion**.
- **Only** a process may be at the same time in the **critical section** of a resource.



Problems with critical sections

20

□ Deadlocks.

▣ Produced with mutual exclusion for multiple resources.

- Process P1 enters critical section for resource A.
- Process P2 enters critical section for resource B.
- Process P1 requests to enter into critical section for resource B.
 - Keeps waiting P2 to exit from critical section for resource B.
- Process P2 requests to enter into critical section for resource A.
 - Keeps waiting P1 to exit from critical section for resource A.

None of them can make progress

Problems with critical sections

21

□ Starvation.


□ A process keeps indefinitely blocked waiting to enter a critical section.

- Process P1 enters critical section for resource A.
- Process P2 requests to enter into critical section for resource A.
 - Keeps waiting for P1 to exit from critical section for resource A.
- Process P3 requests to enter into critical section for resource A
 - Keeps waiting for P1 to exit from critical section for resource A.
 - P2 and P3 are blocked.
- Process P1 leaves critical section for resource A.
- Process P2 enters critical section for resource A.
 - P3 is blocked.
- Process P1 requests to enter into critical section for resource A.
 - Keeps waiting for P2 to exit from critical section for resource A.
 - P2 and P3 are blocked.
- Process P2 leaves critical section for resource A.
 - P3 is blocked
- Process P1 enters into critical section for resource A.
 - P3 is blocked.
- ...

Process P3 never gets to enter into critical section for A.

Starvation

22



Process P3 never gets to enter into critical section for A.

Conditions for mutual exclusion

23

- Only **one process** is allowed to be **at the same time** into the critical section of a resource.
- It must not be possible for a process that requests to enter into a critical section to be **postponed indefinitely**.
- When **no process** is in a critical section, any process requesting to enter that critical section will do so **without delay**.
- No assumptions about **relative speed** of processes or about the **number of processors** must be made.
- A process stays in its critical section for a **finite amount of time**.

Critical section:

Synchronization mechanism

24

- Any mechanism solving the problem of the critical section must provide synchronization between processes.
 - ▣ Each process must request permission to enter into the critical section.
 - ▣ Each process must signal when it is exiting the critical section.

Non-critical code

...

<Entry to critical section>

Critical section code

<Exit from critical section>

...

Non-critical code

Implementation alternatives

25

- ❑ Disable interrupts.
 - ❑ Process would not be interrupted.
 - ❑ Only valid in mono-processor systems.

- ❑ Machine instructions.
 - ❑ Test and set or swap.
 - ❑ Implies busy waiting.
 - ❑ Starvation and deadlock are possible.

- ❑ Other alternative: Operating System support.

Contents

26

- Concurrency.
- Race conditions.
- Mutual exclusion and critical section.
- **Semaphores.**
- Producer-Consumer problem.
- Readers-Writers problem.

Semaphores (Dijkstra)

27

- Process synchronization through a signaling mechanism.
 - Semaphore.

- A semaphore may be seen as an integer variable with three associated operations.
 - ▣ Initialize to a non-negative value.
 - ▣ **semWait**: Decrements the counter
 - If $s < 0 \rightarrow$ Process blocks.
 - ▣ **semSignal**: Increments value of counter.
 - If $s \leq 0 \rightarrow$ Unblocks process.

**Atomic
Operations**

Critical sections and semaphores

28

- A semaphore associated to the critical section of a resource.
- Semaphore initialized to 1.
- Entry to critical section:
semWait.
- Exit from critical section:
semSignal.

Non-critical code

...

semWait(s) ;

Critical section code


semSignal(s) ;

...

Non-critical code

Critical sections and semaphores

29



Contents


30

- Concurrency.
- Race conditions.
- Mutual exclusion and critical section.
- Semaphores.
- **Producer-Consumer problem.**
- Readers-Writers problem.

Producer-Consumer problem

31

- A process **produces** information elements.
- A process **consumes** information elements.
- There is an **intermediate** storage space (buffer).



Infinite buffer

32

Producer


```
for (;;) {  
    x= produce();  
    v[end] = x;  
    end++;  
}
```

Synchronization needed

Consumer

```
for (;;) {  
    while (start==end)  
        {}  
    y=v[start];  
    start++;  
    process(y);  
}
```

**Busy
waiting**



Infinite buffer

33

semaphore s=1

Producer

```
for (;;) {  
    x= produce();  
    semWait(s);  
    v[end] = x;  
    end++;  
    semSignal(s);  
}
```

Consumer

```
for (;;) {  
    while (start==end)  
        {}  
    semWait(s);  
    y=v[start];  
    start++;  
    semSignal(s);  
    process(y);  
}
```

**Busy
waiting**



Infinite buffer

34


`semaphore s=1; semaphore n=0;`

Producer

```
for (;;) {  
    x= produce();  
    semWait(s);  
    v[end] = x;  
    end++;  
    semSignal(s);  
    semSignal(n)  
}
```

Consumer

```
int m;  
for (;;) {  
    semWait(n);  
    semWait(s);  
    y=v[start];  
    start++;  
    semSignal(s);  
    consume(y);  
}
```



Contents

35

- Concurrency.
- Race conditions.
- Mutual exclusion and critical section.
- Semaphores.
- Producer-Consumer problem.
- **Readers-Writers problem.**

Readers-Writers problem

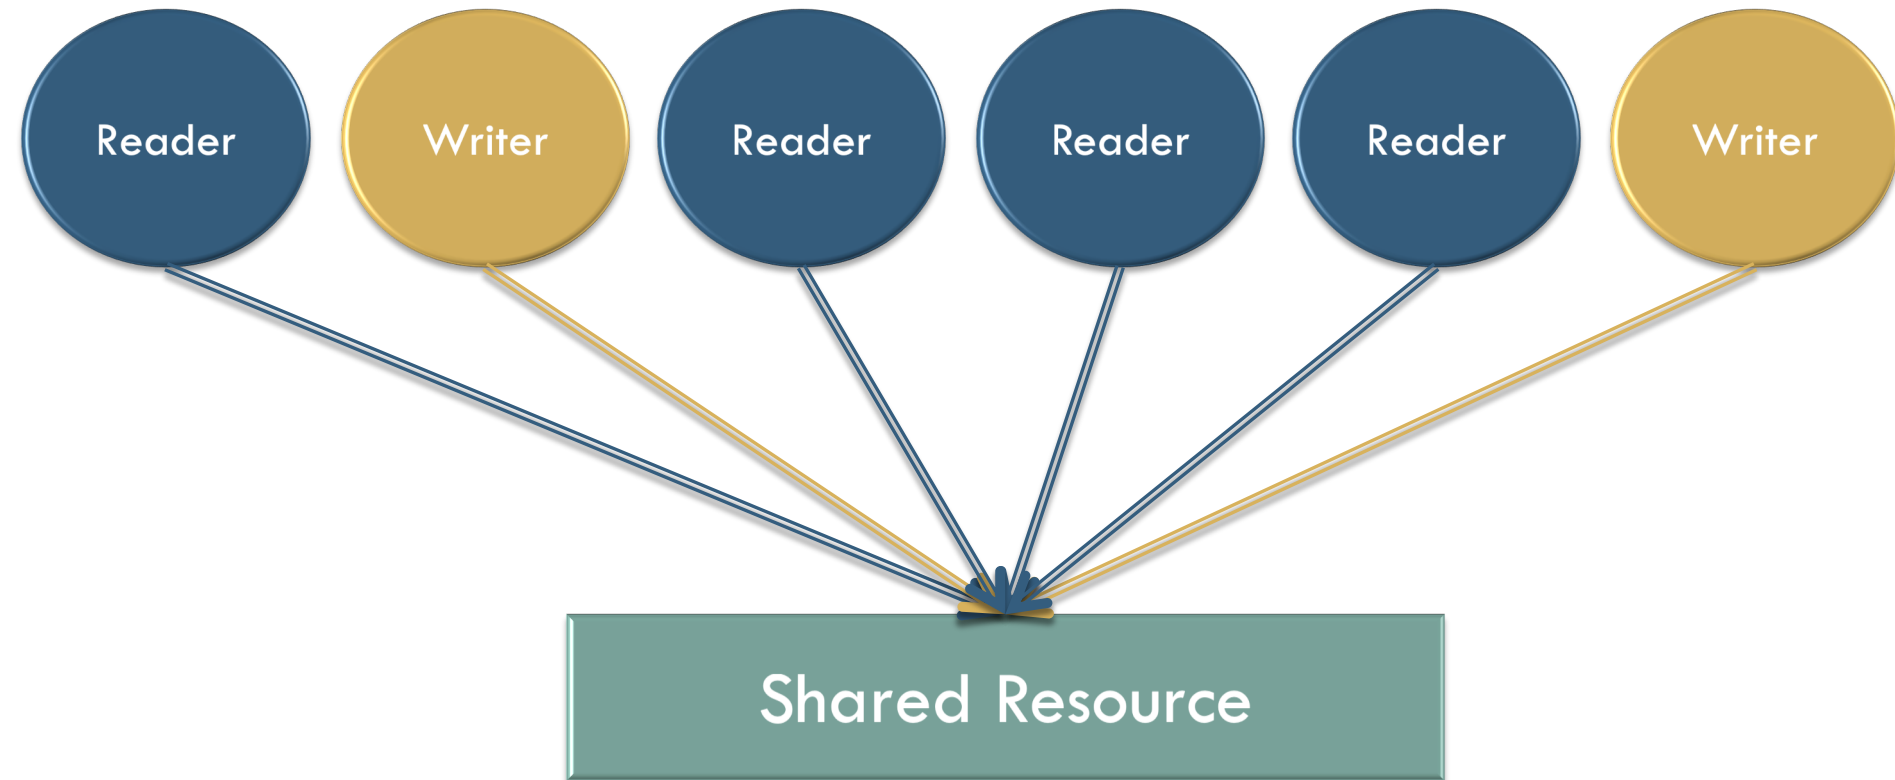
36

- Problem related to a shared area of storage.
 - Multiple processes reading information.
 - Multiple processes writing information.

- **Conditions:**
 - Any number of readers may read from data area concurrently.
 - Only a writer may modify information at once.
 - During a write no reader may perform a request.

Readers-Writers problem

37



Difference with other problems

38

- **Mutual exclusion:**
 - In case of mutual exclusion only a process would be allowed to access information.
 - No concurrency among readers.
- **Producer consumer:**
 - In producer/consumer both processes modify the shared memory area.
- **Goals of additional restrictions:**
 - Provide a more efficient solution.
 - **Reduce contention.**

Management alternatives

39

□ Readers have priority.

- If there is some reader in the critical section other readers may enter.
- A writer may only enter into the critical section if there is no other process already in the critical section.
- Problem: Starvation for writers.

□ Writers have priority.

- When a writer wants to enter into the critical section, no new reader is allowed to enter into the critical section.

Readers have priority

40

```
int nreaders; semaphore readers=1; semaphore writers=1;
```

Reader

```
for(;;) {  
    semWait(readers);  
    nreaders++;  
    if (nreaders==1)  
        semWait(writers);  
    semSignal(readers);
```

```
    perform_read();
```

```
    semWait(readers);  
    nreaders--;  
    if (nreaders==0)  
        semSignal(writers);  
    semSignal(readers);
```

```
}
```

Writer

```
for(;;) {  
    semWait(writers);  
    perform_write();  
    semSignal(writers);  
}
```

**Task: Try to design a solution
for writers have priority**

OPERATING SYSTEMS: PROCESS COMMUNICATION AND SYNCHRONIZATION

Concurrent Processes and problems in communication and synchronization