

# OPERATING SYSTEMS: PROCESS COMMUNICATION AND SYNCHRONIZATION

Threads and communication and synchronization mechanisms

# Content

2

- ❑ **Communication and synchronization.**
- ❑ Semaphores.
- ❑ The readers/writers problem.
  - ❑ Semaphores based solution.
- ❑ Mutex and condition variables.

# Communication mechanisms

3

- **Communication mechanisms** allow for information **transfer** between two processes.
- Files.
- Pipes (pipes, FIFOs).
- **Shared memory variables.**
- Message passing.

# Synchronization mechanisms

4

- Synchronization mechanisms allow to enforce that a process stops its execution until an event happens in another process.
- Concurrent languages constructs (threads).
- Operating system services:
  - ▣ Signals (asynchrony).
  - ▣ Pipes (pipes, FIFOs).
  - ▣ **Semaphores.**
  - ▣ **Mutex and condition variables.**
  - ▣ Message passing.
- Synchronization operations must be **atomic**.

# Content

5

- Communication and synchronization.
- **Semaphores.**
- The readers/writers problem.
  - ▣ Semaphores based solution.
- Mutex and condition variables.

# Semaphores

6

- ❑ Synchronization mechanism.
- ❑ Within the same machine.
- ❑ Object with an associated integer value.
- ❑ Two **atomic** operations.
  - ❑ **wait**
  - ❑ **signal**

# POSIX semaphores

7

- Synchronization mechanism for processes or threads running in the same machine
- POSIX semaphores come in two forms:
  - ▣ **Named semaphores:** can be used by different processes just knowing the name. Does not required shared memory.
  - ▣ **Unnamed semaphores:** can be used only by the process that created it (with threads) or by other using a shared memory region.

```
#include <semaphore.h>
```

```
sem_t * semaphore; //named
```

```
sem_t semaphore; // un-named
```

# POSIX semaphores

8

- **int sem\_init(sem\_t \*sem, int shared, int val);**
  - Initializes unnamed semaphore.
- **int sem\_destroy(sem\_t \*sem);**
  - Destroys unnamed semaphore.
- **sem\_t \*sem\_open(char \*name, int flag, mode\_t mode, int val);**
  - Opens (creates) a named semaphore.
- **int sem\_close(sem\_t \*sem);**
  - Closes a named semaphore.
- **int sem\_unlink(char \*name);**
  - Deletes a named semaphore.
- **int sem\_wait(sem\_t \*sem);**
  - Performs wait operation on a semaphore.
- **int sem\_trywait (sem\_t \*sem);**
  - Try wait. If blocked returns -1 without doing anything.
- **int sem\_post(sem\_t \*sem);**
  - Performs **signal** operation on a semaphore.



## 9

```
sem_post (s); /* critical section exit */
```

- 
- The diagram illustrates the execution of three processes ( $P_0$ ,  $P_1$ ,  $P_2$ ) using a semaphore with an initial value of 1. The vertical axis represents the 'Value of Semaphore (s)', ranging from -2 to 1. Red bars indicate when a process is in its critical section, and black vertical lines indicate when a process is blocked in the semaphore.
- Process  $P_0$ :** Starts with  $s=1$ . It calls `wait(s)`, decreasing  $s$  to 0. It then enters its critical section (red bar). After finishing, it calls `signal(s)`, increasing  $s$  to -1.
  - Process  $P_1$ :** Attempts to call `wait(s)` but finds  $s=-1$  and becomes blocked (black line). When  $P_0$  calls `signal(s)`,  $P_1$  is unblocked,  $s$  increases to 0, and  $P_1$  enters its critical section (red bar). After finishing, it calls `signal(s)`, increasing  $s$  to -1.
  - Process  $P_2$ :** Attempts to call `wait(s)` but finds  $s=-1$  and becomes blocked (black line). When  $P_1$  calls `signal(s)`,  $P_2$  is unblocked,  $s$  increases to 0, and  $P_2$  enters its critical section (red bar). After finishing, it calls `signal(s)`, increasing  $s$  to 1.
- Legend:
- Red bar: Executing code in critical section
  - Black line: Process blocked in semaphore

# Operations on semaphores

10

```
sem_wait(s) {  
    s = s - 1;  
    if (s < 0) {  
        <Block process>  
    }  
}
```

```
sem_post (s) {  
    s = s + 1;  
    if (s <= 0) {  
        <Unblock a blocked process by wait operation>  
    }  
}
```

# Unnamed semaphores:

## Producer consumer with semaphores


11

```
#define MAX_BUFFER    1024  /* buffer size*/
#define DATA_SIZE 100000  /* number of data items to produce */

sem_t elements;          /* elements in buffer*/
sem_t holes;             /* holes in buffer*/
int buffer[MAX_BUFFER];  /* common buffer*/

void main(void)
{
    pthread_t th1, th2; /* threads identifiers*/

    /* Initialize semaphores*/
    sem_init(&elements, 0, 0);
    sem_init(&holes, 0, MAX_BUFFER);
```



# Unnamed semaphores:

## Producer consumer with semaphores

12

```
/* create threads*/  
pthread_create(&th1, NULL, producer, NULL);  
pthread_create(&th2, NULL, consumer, NULL);  
  
/* wait for termination*/  
pthread_join(th1, NULL);  
pthread_join(th2, NULL);  
  
sem_destroy(&holes);  
sem_destroy(&elements);  
  
exit(0);  
}
```

# Unnamed semaphores:

## Producer & Consumer threads

13

```
void producer() { /* Producer code*/
    int pos = 0; /* position in buffer*/
    int data; /* data to be produced */
    int i;

    for(i=0; i < DATA_SIZE; i++ ) {
        data = i; /* produce data*/

        sem_wait(&holes); /* Reduce holes by 1*/
        buffer[pos] = i;
        pos = (pos + 1) % MAX_BUFFER;
        sem_post(&elements); /* one element
                               more*/
    }
    pthread_exit(0);
}
```

```
void consumer() { /* Consumer code */
    int pos = 0;
    int data;
    int i;

    for(i=0; i < DATA_SIZE; i++ ) {
        sem_wait(&elements); /* one element less*/
        data= buffer[pos];
        pos = (pos + 1) % MAX_BUFFER;
        sem_post(&huecos); /* one hole more*/

        store_data(data); /* consume data */
    }
    pthread_exit(0);
}
```

# Content

14

- Communication and synchronization.
- Semaphores.
- **The readers/writers problem.**
  - ▣ Semaphores based solution.
- Mutex and condition variables.

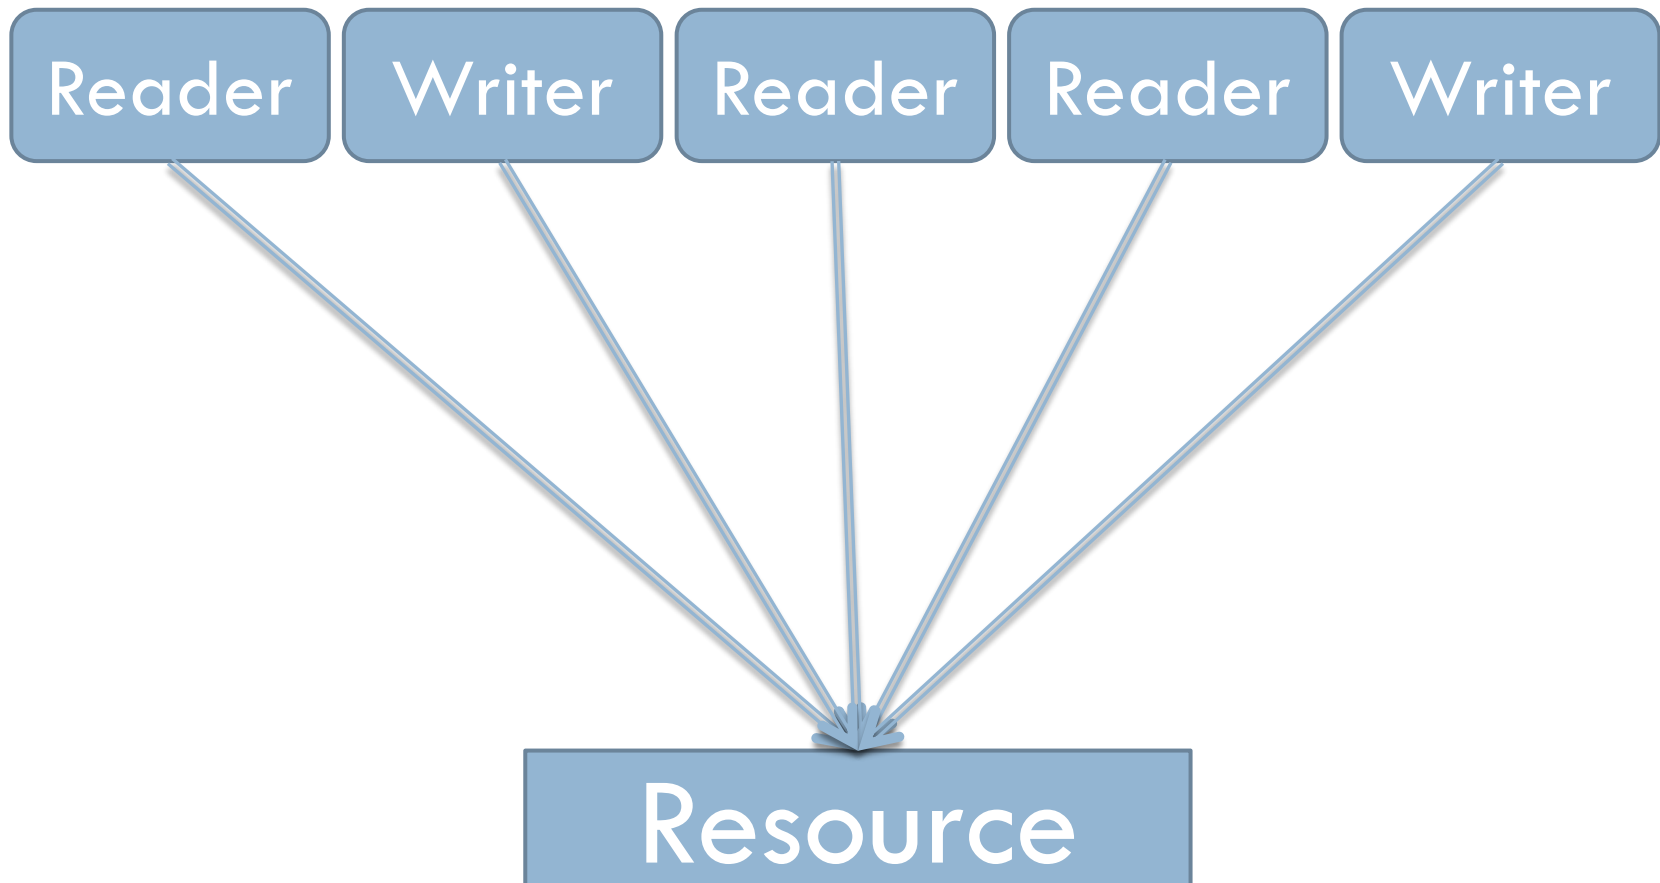
# Readers-writes problem

15

- Problem when using a shared storage area.
  - Multiple processes reading information.
  - Multiple processes writing information.
  
- **Conditions:**
  - ▣ Any number of readers may read from the data area concurrently.
    - ▣ Read concurrency allowed.
  - ▣ Only one writer may modify information at once.
    - ▣ No concurrent writes allowed.
  - ▣ During a write no reader may perform a query.
    - ▣ No read/write concurrency allowed.

# Readers-writers problem

16





# Differences with other problems

17

- **Mutual exclusion**

- Mutual exclusion would only allow a process to access at the same time to information.
- No concurrency allowed between readers.
- Higher contention.

- **Producer/Consumer:**

- In producer/consumer both processes modify the shared data area.

- Goals of additional constraints:

- **Provide a more efficient solution.**

# Management alternatives

18

## □ **Readers have priority.**

- If there is some reader in critical section other readers may enter.
- A writer can only enter the critical section if there is no process there.
- **Problem:** Starvation for writers.

## □ **Writers have priority.**

- When a writer wishes to access to the critical section no more readers are admitted.

# Readers have priority

19

```
int nreaders; semaphore rd=1; semaphore wr=1;
```

## Lector

```
for(;;) {  
    semWait(rd);  
    nreaders++;  
    if (nreaders==1)  
        semWait(wr);  
    semSignal(rd);  
  
    perform_read();  
  
    semWait(rd);  
    nreaders--;  
    if (nreaders==0)  
        semSignal(wr);  
    semSignal(rd);  
}
```

## Escritor

```
for(;;) {  
    semWait(wr);  
    perform_write();  
    semSignal(wr);  
}
```

# Unnamed semaphores:

## Readers-writers with semaphores

20

```
int data = 5;    /* resource*/
int nreaders = 0; /* number of readers */
sem_t sem_rd;    /* controls access to nreaders*/
sem_t mutex;     /* controls access to data*/

void main(void) {
    pthread_t th1, th2, th3, th4;

    sem_init(&mutex, 0, 1);
    sem_init(&sem_rd, 0, 1);

    pthread_create(&th1, NULL, reader, NULL);
    pthread_create(&th2, NULL, writer, NULL);
    pthread_create(&th3, NULL, reader, NULL);
    pthread_create(&th4, NULL, writer, NULL);
}
```

# Unnamed semaphores:

## Readers-writers with semaphores

21

```
pthread_join(th1, NULL);  
pthread_join(th2, NULL);  
pthread_join(th3, NULL);  
pthread_join(th4, NULL);  
  
/* cerrar todos los semaforos */  
sem_destroy(&mutex);  
sem_destroy(&sem_rd);  
  
exit(0);  
}
```

# Unnamed semaphores:

## Readers & writers threads

22

```
void reader() { /* reader code */  
    sem_wait(&sem_rd);  
    nreaders = nreaders + 1;  
    if (nreaders == 1) sem_wait(&mutex);  
    sem_post(&sem_rd);  
  
    printf("`%d\n", data); /* read data and print  
        it*/  
  
    sem_wait(&sem_rd);  
    n_readers = n_readers - 1;  
    if (n_readers == 0) sem_post(&mutex);  
    sem_post(&sem_rd);  
    pthread_exit(0);  
}
```

```
void writer() { /* writer code */  
    sem_wait(&mutex);  
    data = data + 2; /* modify resource  
        */  
    sem_post(&mutex);  
  
    pthread_exit(0);  
}
```

# Named semaphores:

## Names

23

- To synchronize processes without using shared memory.
- Name: string similar to a file name.
  - ▣ If name is relative, can only be accessed by the creator and its children.
  - ▣ If name is absolute (starts with “/”), the semaphore can be shared by any process knowing the name and having permissions.
- Standard way to create semaphores to be used by parent and children
  - ▣ Unnamed semaphores not valid, as parent and children DO NOT share memory.

# Named semaphores:

## Creation and use

24

### □ To create it:

```
sem_t *sem_open(char *name, int flag, mode_t mode, int val);
```

- Flag = O\_CREAT create it.
- Flag: O\_CREAT | O\_EXECL. Create, but returns -1 if exist.
- Mode: access permissions;
- Val: initial value of the semaphore ( $\geq 0$ );

### □ To use it:

```
sem_t *sem_open(char *name, int flag);
```

- With flag 0. It does not exist returns -1.

### □ Important:

- All processes must know the “name” and use the same.



# Named semaphores:

## Readers - Writers

25

```
int dato = 5;    /* recurso */
int n_lectores = 0; /* num lectores */
sem_t *sem_lec; sem_t *mutex;

int main (int argc, char *argv[]) {

    int i, n= 5; pid_t pid;
    /* Named semaphore*/
    if((mutex=sem_open("/tmp/sem_1", O_CREAT, 0644,
1))==(sem_t *)-1)
        { perror("It is possible to create"); exit(1); }
    if((sem_lec=sem_open("/tmp/sem_2", O_CREAT,
0644, 1))==(sem_t *)-1)
        { perror(" It is possible to create "); exit(1); }
```

```
/* Crea los procesos */
for (i = 1; i< atoi(argv[1]); ++i){
    pid = fork();
    if (pid ==-1)
        { perror("No se puede crear el proceso");
          exit(-1);}
    else if(pid==0) { /child
        reader(getpid()); break;
    }
    writer(pid); /* parent */
}

sem_close(mutex); sem_close(sem_lec);
sem_unlink("/tmp/sem_1");
sem_unlink("/tmp/sem_2");
```

# Named semaphores:

## Readers & writers processes

26

```
void reader (int pid) {  
    sem_wait(sem_lec);  
    n_lectores = n_lectores + 1;  
    if (n_lectores == 1)  
        sem_wait(mutex);  
    sem_post(sem_lec);  
    printf(" lector %d  dato: %d\n", pid, dato); /* leer dato */  
    sem_wait(sem_lec);  
    n_lectores = n_lectores - 1;  
    if (n_lectores == 0)  
        sem_post(mutex);  
    sem_post(sem_lec);  
}
```

```
void writer (int pid) {  
    sem_wait(mutex);  
    dato = dato + 2; /* modificar  
    el recurso */  
    printf("escritor %d  dato: %d\n",  
    pid, dato); /* leer dato */  
    sem_post(mutex);  
}
```

# Content

27

- Communication and synchronization.
- Semaphores.
- The readers/writers problem.
  - ▣ Semaphores based solution.
- **Mutex and condition variables.**

# Mutex and condition variables

28


- A **mutex** is a synchronization mechanism for threads.
- It is a binary semaphore with to **atomic** operations.
  - **lock(m)** Try to block the mutex. If the mutex is already blocked the calling thread is suspended.
  - **unlock(m)** Unblocks the mutex. If there are processes blocked in the mutex one is unblocked.

# Critical sections with mutexes

29

```
lock(m);    /* entry into critical section */  
<critical section>  
unlock(s); /* exit from critical section */
```

- **unlock** operation must be performed by the thread that performed **lock**



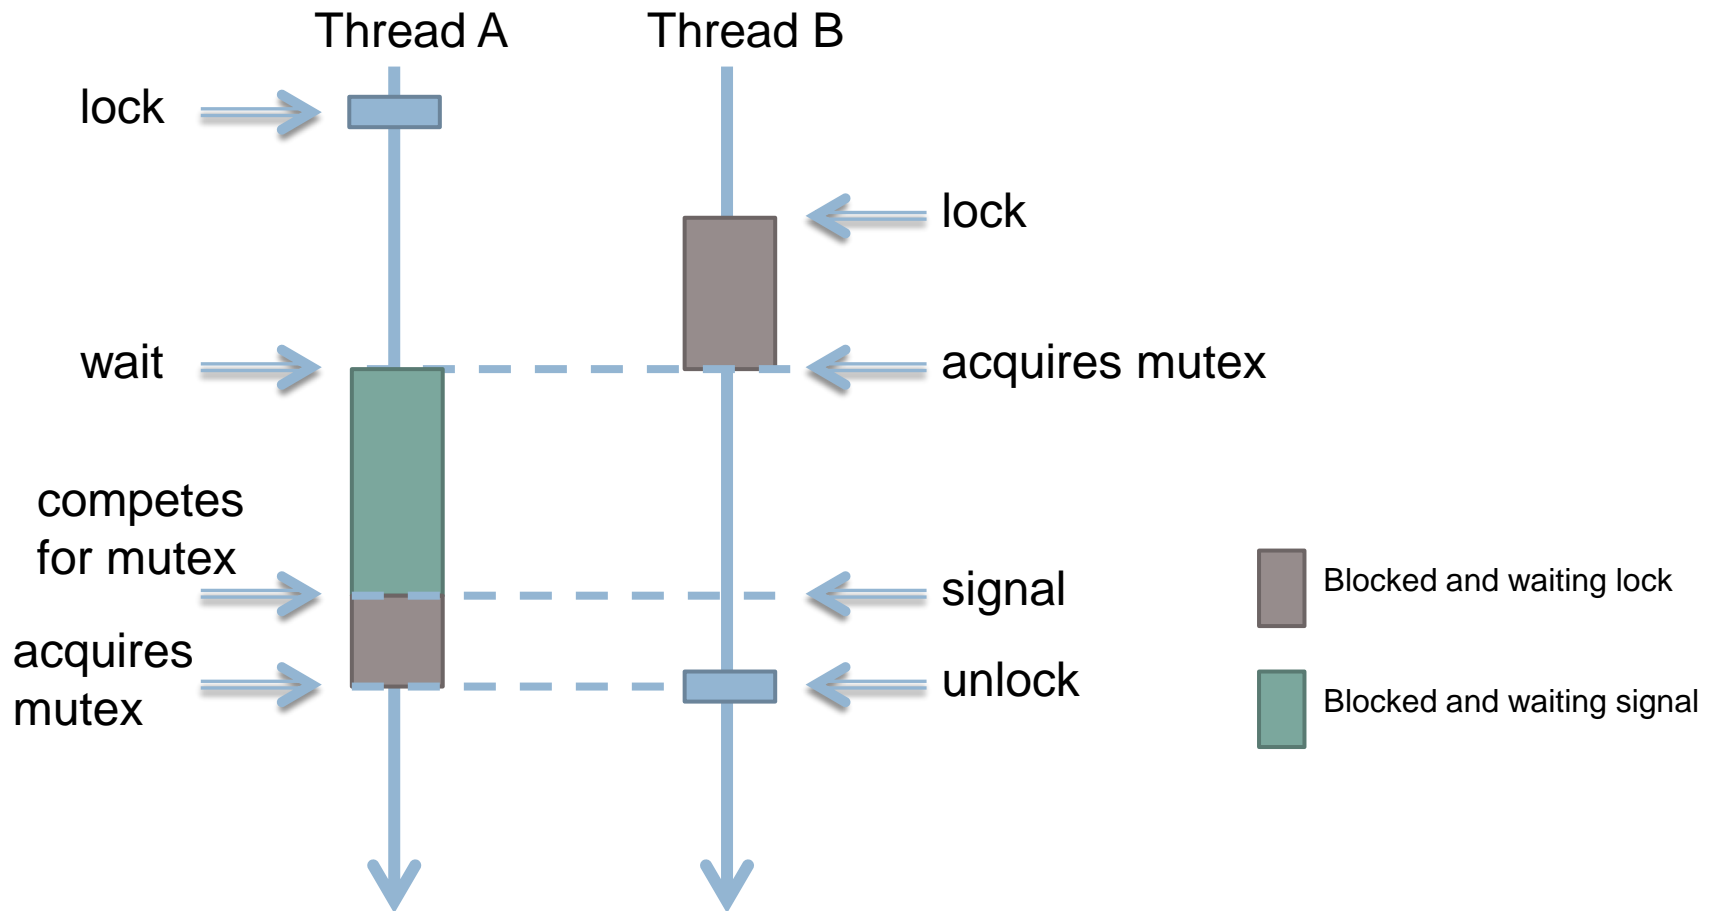
# Condition variables

30

- Synchronization variables associated to a mutex.
- Two **atomic** operations:
  - **wait** Blocks running thread and releases mutex.
  - **signal** Unblocks one or more threads suspended in the condition variable. Unblocked threads contend for acquiring the mutex again.
- It is convenient to run them in a **lock/unlock** block.

# Condition variables

31



# Using mutexes and condition variables

32

## Thread A

```
lock(mutex); /* access to resource */  
<check data structures>  
while (resource is busy) {  
    wait (condition, mutex);  
}  
<mark resource as busy>  
unlock (mutex);
```

## Thread B

```
lock (mutex); /* access to resource*/  
<mark resource as free>  
signal (condition, mutex);  
unlock (mutex);
```

Important to use while



# POSIX services

33

```
int pthread_mutex_init(pthread_mutex_t *mutex, pthread_mutexattr_t * attr);
```

- ▣ Initialize mutex.

```
int pthread_mutex_destroy(pthread_mutex_t *mutex) ;
```

- ▣ Destroy mutex.

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

- ▣ Try to get access to mutex.
- ▣ Blocks thread if mutex is already acquired by other thread.

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- ▣ Unblock mutex.

# POSIX services

34

```
int pthread_cond_init(pthread_cond_t*cond, pthread_condattr_t*attr);
```

- ▣ Initialize a condition variable.

```
int pthread_cond_destroy(pthread_cond_t *cond);
```

- ▣ Destroy a condition variable.

# POSIX services

35

```
int pthread_cond_signal(pthread_cond_t *cond);
```

- ▣ Unblocks one or more threads that are suspended in the condition variable **cond**.
- ▣ Has no effect if there is no thread waiting (difference with semaphores).

```
int pthread_cond_broadcast(pthread_cond_t *cond);
```

- ▣ All blocked threads in condition variable **cond** are unblocked.
- ▣ Has no effect if there is no thread waiting.

```
int pthread_cond_wait(pthread_cond_t*cond, pthread_mutex_t*mutex);
```

- ▣ Suspend thread until another thread signals condition variable **cond**.
- ▣ Automatically releases the **mutex**. When thread is unblocked it contends again for the **mutex**.

# Producer-Consumer with mutexes

36

```
#define MAX_BUFFER    1024    /* size of buffer*/
#define DATA_SIZE  100000    /* number of data items to be produced*/

pthread_mutex_t mutex; /* mutex to access shared buffer */
pthread_cond_t non_full; /* can we add more elements? */
pthread_cond_t non_empty; /* can we remove elements? */
int n_elements;    /* number of elements in buffer */
int buffer[MAX_BUFFER]; /* common buffer */

int main() {
    pthread_t th1, th2;

    pthread_mutex_init(&mutex, NULL);
    pthread_cond_init(&non_full, NULL);
    pthread_cond_init(&non_empty, NULL);
```

# Producer-Consumer with mutexes

37

```
pthread_create(&th1, NULL, producer, NULL);
```

```
pthread_create(&th2, NULL, consumer, NULL);
```

```
pthread_join(th1, NULL);
```

```
pthread_join(th2, NULL);
```

```
pthread_mutex_destroy(&mutex);
```

```
pthread_cond_destroy(&non_full);
```

```
pthread_cond_destroy(&non_empty);
```

```
exit(0);
```

```
}
```

# Producer

38

```
void producer() { /* Producer code */
    int data, i ,pos = 0;
    for(i=0; i < DATA_SIZE; i++ ) {
        data= i;    /* generate data */
        pthread_mutex_lock(&mutex);    /* access to buffer*/
        while (n_elements == MAX_BUFFER) /* when buffer is full*/
            pthread_cond_wait(&non_full, &mutex);
        buffer[pos] = data;
        pos = (pos + 1) % MAX_BUFFER;
        n_elements ++;
        pthread_cond_signal(&non_empty); /* buffer is not empty */
        pthread_mutex_unlock(&mutex);
    }
    pthread_exit(0);
}
```

# Consumer

39

```
void consumer() { /* consumer code */
    int dato, i ,pos = 0;
    for(i=0; i < DATA_SIZE; i++ ) {
        pthread_mutex_lock(&mutex); /* access to buffer */
        while (n_elements == 0) /* when buffer empty */
            pthread_cond_wait(&non_empty, &mutex);
        dato = buffer[pos];
        pos = (pos + 1) % MAX_BUFFER;
        n_elements --;
        pthread_cond_signal(&non_full); /* buffer is not full */
        pthread_mutex_unlock(&mutex);
        printf("Consumed %d \n", dato); /* Use data*/
    }
    pthread_exit(0);
}
```

# Readers writers with mutex

40

```
int data= 5;          /* resource*/
int nreaders = 0;     /* number of readers */
pthread_mutex_t data_mutex;    /* Control access to data*/
pthread_mutex_t mutex_rd; /* Controls access to nreaders */

main(int argc, char *argv[]) {
    pthread_t th1, th2, th3, th4;

    pthread_mutex_init(&data_mutex, NULL);
    pthread_mutex_init(&mutex_rd, NULL);

    pthread_create(&th1, NULL, Lector, NULL);
    pthread_create(&th2, NULL, Escritor, NULL);
    pthread_create(&th3, NULL, Lector, NULL);
    pthread_create(&th4, NULL, Escritor, NULL);
```



# Lectores-escriptores con mutex

41

```
pthread_join(th1, NULL);  
pthread_join(th2, NULL);  
pthread_join(th3, NULL);  
pthread_join(th4, NULL);  
  
pthread_mutex_destroy(&data_mutex);  
pthread_mutex_destroy(&mutex_rd);  
  
exit(0);  
}
```

# Writer

42

```
void writer() { /* writer code*/  
    pthread_mutex_lock(&data_mutex);  
    dato = dato + 2; /* modify resource*/  
    pthread_mutex_unlock(&data_mutex);  
    pthread_exit(0);  
}
```

# Reader

43

```
void reader() { /* codigo del lector */  
    pthread_mutex_lock(&mutex_rd);  
    nreaders++;  
    if (nreaders == 1) pthread_mutex_lock(&data_mutex);  
    pthread_mutex_unlock(&mutex_rd);  
  
    printf("%d\n", data); /* read data and print it*/  
  
    pthread_mutex_lock(&mutex_rd);  
    nreaders--;  
    if (nreaders == 0) pthread_mutex_unlock(&data_mutex);  
    pthread_mutex_unlock(&mutex_rd);  
  
    pthread_exit(0);  
}
```

# OPERATING SYSTEMS: PROCESS COMMUNICATION AND SYNCHRONIZATION

Threads and communication and synchronization mechanisms