

OPERATING SYSTEMS: PROCESS COMMUNICATION AND SYNCHRONIZATION

Developing concurrent servers

Contents

2

- ❑ **Request servers.**
- ❑ Process-based solution.
- ❑ Threads on demand solution.
- ❑ Thread pool solution.


Request servers

3

- Request servers are developed in many contexts:
 - Web server.
 - Database server.
 - Applications server.
 - File exchange servers.
 - Instant messaging applications.
 - ...


Server

4



Problem: Request server

5

- A server receives requests that must be processed.
 - Generic server structure:
 - Request **reception**.
 - Each request requires a certain amount of time in input/output operations to be received.
 - Request **processing**.
 - A certain amount of time for CPU processing.
 - Reply **sending**.
 - A certain amount of input/output operations to perform reply.
- 

A library for testing

6

- To be able to evaluate solutions today we will use a simple library as a baseline example.

- Ideas:
 - ▣ Simulate requests reception.
 - ▣ Simulate request processing.
 - ▣ Simulate reply sending.

Base library

7

```
#ifndef REQUEST_H
#define REQUEST_H

struct request {
    long id;
    /* Rest of fields */
    int kind;
    char url[80];
    /* ... */
};

typedef struct request request_t;

void get_request (request_t * r);
void reply_request (request_t * r);

#endif
```

Receiving requests

8

```
static long reqid = 0;

void receive_request(request_t * r) {
    int delay;
    fprintf(stderr, "Receiving request\n");
    r->id = reqid++;

    /* Simulate I/O time*/
    delay = rand() % 5;
    sleep(delay);

    fprintf(stderr, "Request %d received after %d seconds\n", r->id, delay);
}
```


Receiving requests

9

```
static long reqid = 0;
```

```
void receive_request(request_t * r) {  
    int delay;  
    fprintf(stderr, "Receiving request\n");  
    r->id = reqid++;
```

```
    /* Simulate I/O time*/
```

```
    delay = rand() % 5;
```

```
    sleep(delay);
```

```
    fprintf(stderr, "Request %d received after %d seconds\n", r->id, delay);
```

```
}
```

In real code here some blocking call would be used to receive the request (e.g. from network)

Sending requests

10

```
void reply_request (request_t * r) {  
    int delay, i;  
    double x;  
    fprintf(stderr, "Sending reply %d\n", r->id);  
  
    /* Simulate processing time */  
    for (i=0;i<1000000;i++) { x = 2.0 * i; }  
  
    /* Simulate I/O time */  
    delay = rand() % 20;  
    sleep(delay);  
  
    fprintf(stderr, "Request %d sent after %d seconds\n", r->id, delay);  
}
```

Sending requests

11

```
void reply_request (request_t * r) {  
    int delay, i;  
    double x;  
    fprintf(stderr, "Sending reply %d\n", r->id);
```

```
    /* Simulate processing time */  
    for (i=0;i<1000000;i++) { x = 2.0 * i; }
```

Request
processing would
go here

```
    /* Simulate I/O time */  
    delay = rand() % 20;  
    sleep(delay);
```

Some blocking call would be
used here to reply the
request.

```
    fprintf(stderr, "Request %d sent after %d seconds\n", r->id, delay);  
}
```

A first try

12

- Run forever the following sequence:
 - ▣ Receive a request.
 - ▣ Reply the request.

```
#include "request.h"

int main() {
    request_t r;

    for (;;) {
        receive_request(&r);
        reply_request(&r);
    }

    return 0;
}
```

Problems

13

- Request arrival.
 - ▣ If **two requests** arrive **at the same time** ...
 - ▣ If **a request** arrives while **one is being processed** ...

- Resource utilization.
 - ▣ How will CPU utilization be?

Measuring initial solution

14

```
#include "request.h"
```

```
#include <stdio.h>
```

```
#include <time.h>
```

```
int main() {
```

```
    int i;
```

```
    const int MAX_REQUESTS = 5;
```

```
    time_t t1,t2;
```

```
    double diff;
```

```
    request_t r;
```

```
t1 = time(NULL);
```

```
for (i=0;i<MAX_REQUESTS;i++) {
```

```
    receive_request(&r);
```

```
    reply_request(&r);
```

```
}
```

```
t2 = time(NULL);
```

```
diff = difftime(t2,t1);
```

```
printf("Time: %lf\n",diff);
```

```
return 0;
```

```
}
```

Execution

15

\$ time app/seq

Receiving request

Request 0 received after 3 seconds

Sending reply 0

Request 0 sent after 6 seconds (computed 999999.000000)

Receiving request

Request 1 received after 2 seconds

Sending reply 1

Request 1 sent after 15 seconds (computed 999999.000000)

Receiving request

Request 2 received after 3 seconds

Sending reply 2

Request 2 sent after 15 seconds (computed 999999.000000)

Receiving request

Request 3 received after 1 seconds

Sending reply 3

Request 3 sent after 12 seconds (computed 999999.000000)

Receiving request

Request 4 received after 4 seconds

Sending reply 4

Request 4 sent after 1 seconds (computed 999999.000000)

Time: 62.000000

real 1m2.009s

user 0m0.008s

sys 0m0.000s

Evaluation

16

Normal	Processes	Thread per request	<i>Thread Pool</i>
62 seg.			

Contents

17

- Request servers.
- **Process-based solution.**
- Threads on demand solution.
- Thread pool solution.

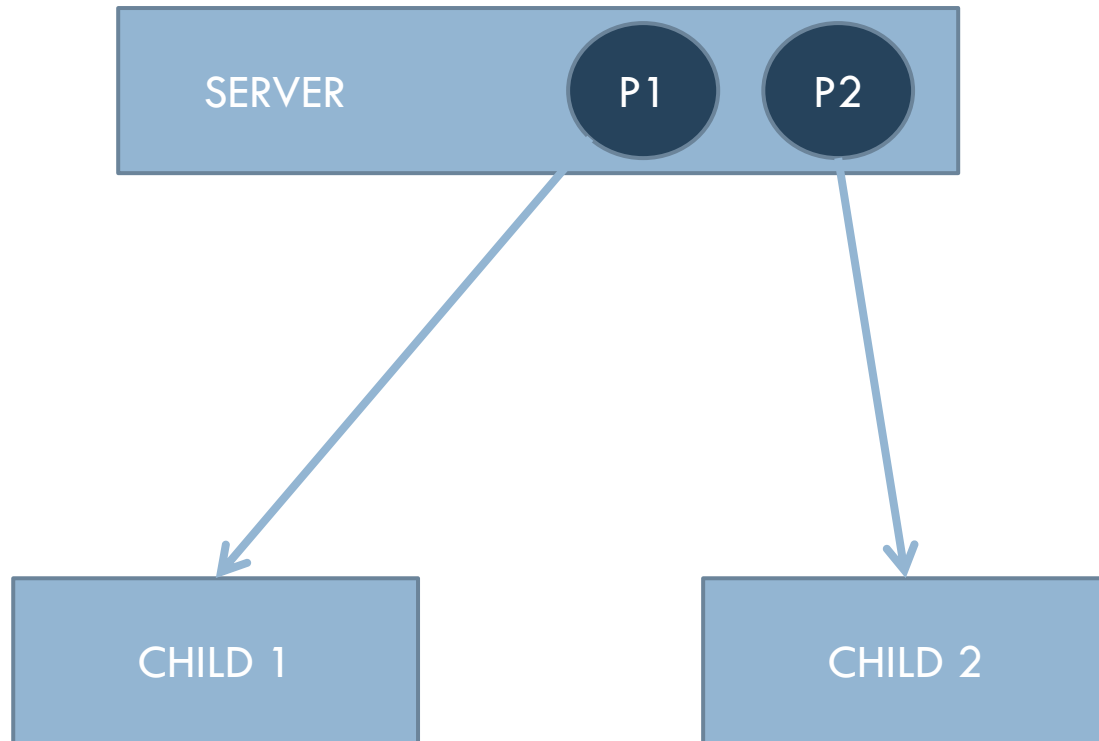
Main idea

18

- Each time a request arrives a new child process is created:
- **Child process** performs **request processing**.
- **Parent process** waits for **next request**.

Process-based server

19



Implementation (1 / 3)

20

```
#include "request.h"  
#include <stdio.h>  
#include <time.h>  
#include <sys/wait.h>
```

```
int main() {  
    const int MAX_REQUESTS = 5;  
    int i;  
    time_t t1,t2;  
    request_t r;  
    int pid, nchildren=0;  
  
    t1 = time(NULL);
```

Implementation (2/3)

21

```
for (i=0;i<MAX_REQUESTS; ++i) {  
    receive_request(&r);  
  
    do {  
        fprintf(stderr, "Checking children\n");  
        pid = waitpid(-1, NULL, WNOHANG);  
        if (pid>0) { nchildren--; }  
    } while (pid > 0);  
  
    pid = fork();  
    if (pid<0) { perror("Error creating child"); }  
    if (pid==0) { reply_request(&r); exit(0); } /* CHILD */  
    if (pid!=0) { nchildren++; }                /* PARENT */  
}
```

Implementation (3/3)

22

```
fprintf(stderr, "Checking %d children\n", nchildren);
while (nchildren>0) {
    pid = waitpid(-1, NULL, 0); /* Blocking wait */
    if (pid>0) { nchildren--; }
};

t2 = time(NULL);

double diff = difftime(t2,t1);
printf("Time: %lf\n",diff);


return 0;
}
```

Execution

23

\$ time app/procsrv

Receiving request
Request 0 received after 3 seconds
Checking children
Receiving request
Sending reply 0
Request 1 received after 1 seconds
Checking children
Receiving request
Sending reply 1
Request 2 received after 2 seconds
Checking children
Receiving request
Request 3 received after 0 seconds
Checking children
Sending reply 2
Receiving request
Sending reply 3



Request 4 received after 3 seconds
Checking children
Checking 5 children
Sending reply 4
Request 0 sent after 6 seconds (computed 999999.000000)
Request 3 sent after 13 seconds (computed 999999.000000)
Request 2 sent after 15 seconds (computed 999999.000000)
Request 1 sent after 17 seconds (computed 999999.000000)
Request 4 sent after 15 seconds (computed 999999.000000)

Time: 24.000000

real 0m24.004s
user 0m0.001s
sys 0m0.012s

Evaluation

24

Normal	Processes	Thread per request	<i>Thread Pool</i>
62 seg.	24 seg.		

Problems

25

- A new process must be started (**fork**) for each incoming request.
- A process must be terminated (**exit**) for each served request.
- **Excessive** system **resources utilization**.
- No **admission control**.
 - Problems with **Quality of Service**.

Solutions with threads

26

□ On-demand threads.

- ▣ Each time a request is received a new thread is created.

□ Thread pool.

- ▣ A fixed number of pre-created threads.
- ▣ Each time a request is received a free thread is assigned to serve the request.
 - Communication using a **request queue**.

Contents

27

- Request servers.
- Process-based solution.
- **Threads on demand solution.**
- Thread pool solution.

On-demand

28

- A receiver thread is in charge of receiving the requests.
- Each time a request arrives a new thread is created, and a **copy** of the request is passed to the new created thread.
 - It must be a copy of the request as the original request could be modified.

Implementation

29

```
#include "request.h"
#include <stdio.h>
#include <time.h>
#include <pthread.h>
#include <semaphore.h>
```

```
sem_t snchildren;
```

```
int main() {
    time_t t1, t2;
    double diff;
    pthread_t thr;
```



```
t1 = time(NULL);
sem_init(&snchildren, 0, 0);
pthread_create(&thr, NULL,
               receiver, NULL);
pthread_join(thr, NULL);
sem_destroy(&snchildren);
t2 = time(NULL);

diff = difftime(t2, t1);
printf("Time: %f\n", diff);

return 0;
}
```



Implementation: receiver

30

```
void * receiver(void * param) {
    const int MAX_REQUESTS = 5; int i, nservice = 0;
    request_t r;
    pthread_t th_child;

    for (i=0;i<MAX_REQUESTS;i++) {
        receive_request(&r);
        nservice++;
        pthread_create(&th_child, NULL, service, &r);
        pthread_detach(th_child);
    }

    while (nservice>0) {
        fprintf(stderr, "Performing wait\n");
        sem_wait(&snchildren);
        nservice--;
        fprintf(stderr, "Exiting from wait\n");
    }

    pthread_exit(0);
    return NULL;
}
```

Implementation: service

31

```
void * service(void * r) {  
    request_t req;  
  
    copy_request(&req,(request_t*)r);  
  
    fprintf(stderr, "Starting service\n");  
    reply_request(&req);  
    sem_post(&snchildren);  
  
    fprintf(stderr, "Finishing service\n");  
    pthread_exit(0);  
    return NULL;  
}
```

Execution

32

\$ app/mtsrv

Receiving request

Request 0 received after 3 seconds

Receiving request

Starting service

Sending reply 1

Request 1 received after 1 seconds

Receiving request

Request 2 received after 0 seconds

Starting service

Sending reply 2

Receiving request

Starting service

Sending reply 3

Request 3 received after 3 seconds

Receiving request

Starting service

Sending reply 4

Request 4 received after 2 seconds

Performing wait

Starting service

Sending reply 4

Request 3 sent after 6 seconds (computed 999999.000000)

Finishing service for 3

Exiting from wait

Performing wait

Request 4 sent after 1 seconds (computed 999999.000000)

Finishing service for 4

Exiting from wait

Performing wait

Execution

33

Request 3 sent after 6 seconds (computed 999999.000000)

Finishing service for 3

Exiting from wait

Performing wait

Request 4 sent after 1 seconds (computed 999999.000000)

Finishing service for 4

Exiting from wait

Performing wait

Request 4 sent after 9 seconds (computed 999999.000000)

Finishing service for 4

Exiting from wait

Performing wait

Request 2 sent after 15 seconds (computed 999999.000000)

Finishing service for 2

Exiting from wait

Performing wait

Request 1 sent after 17 seconds (computed 999999.000000)

Finishing service for 1

Exiting from wait

Time: 20.000000

real 0m20.004s

user 0m0.013s

sys 0m0.000s

Question

34

□ Can a race condition happen?

```
jdgarci@gavilan:~/ejcpp/server/build$ valgrind --tool=helgrind app/mtsrv
```

```
==7970== Helgrind, a thread error detector
```

```
==7970== Copyright (C) 2007-2013, and GNU GPL'd, by OpenWorks LLP et al.
```

```
==7970== Using Valgrind-3.10.0.SVN and LibVEX; rerun with -h for copyright info
```

```
==7970== Command: app/mtsrv
```

```
==7970==
```

```
....
```

```
==7970== Possible data race during write of size 4 at 0x5C1FE38 by thread #2
```

```
==7970== Locks held: none
```

```
==7970== at 0x400F9B: receive_request (in /home/jdgarci/ejcpp/server/build/app/mtsrv)
```

```
==7970== by 0x400ED3: receiver (in /home/jdgarci/ejcpp/server/build/app/mtsrv)
```

```
==7970== by 0x4C30E26: mythread_wrapper (hg_intercepts.c:233)
```

```
==7970== by 0x4E45181: start_thread (pthread_create.c:312)
```

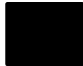
```
==7970== by 0x515547C: clone (clone.S:111)
```

Race condition

35

- Can a race condition happen?

```
void * receiver (void * param)
```


```
request_t r;   
  
receive_request(&r);  
nservice++;  
pthread_create(&child, NULL, servicio, &r);  
  
receive_request(&r);  
nservice++;  
pthread_create(&child, NULL, servicio, &r);  
  
...
```

Race condition

36

- Can a race condition happen?

```
void * receiver(void * param)
```

```
request_t r;   
receive_request(&r);  
nservice++;  
pthread_create(&child, NULL, service, &r);  
  
receive_request(&r);  
nservice++;  
pthread_create(&child, NULL, service, &r);  
...
```

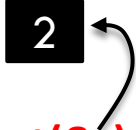
Race condition

37

- Can a race condition happen?

```
void * receiver(void * param)
```

```
request_t r; 2  
receive_request(&r);  
nservice++;  
pthread_create(&child, NULL, service, &r);  
  
receive_request(&r);  
nservice++;  
pthread_create(&child, NULL, service, &r);  
...
```




Race condition

38

- Can a race condition happen?

```
void * receiver(void * param)
```

```
request_t r; 2  
  
receive_request(&r);  
nservice++;  
pthread_create(&child, NULL, service, &r);  
  
receive_request(&r);  
nservice++;  
pthread_create(&child, NULL, service, &r);  
  
...
```



Race condition

39

□ Can a race condition happen?

```
void * receiver(void * param)
```

```
request_t r; 2  
  
receive_request(&r);  
nservice++;  
pthread_create(&hijo, NULL, service, &r);  
  
receive_request(&r);  
nservice++;  
pthread_create(&hijo, NULL, service, &r);  
  
...
```

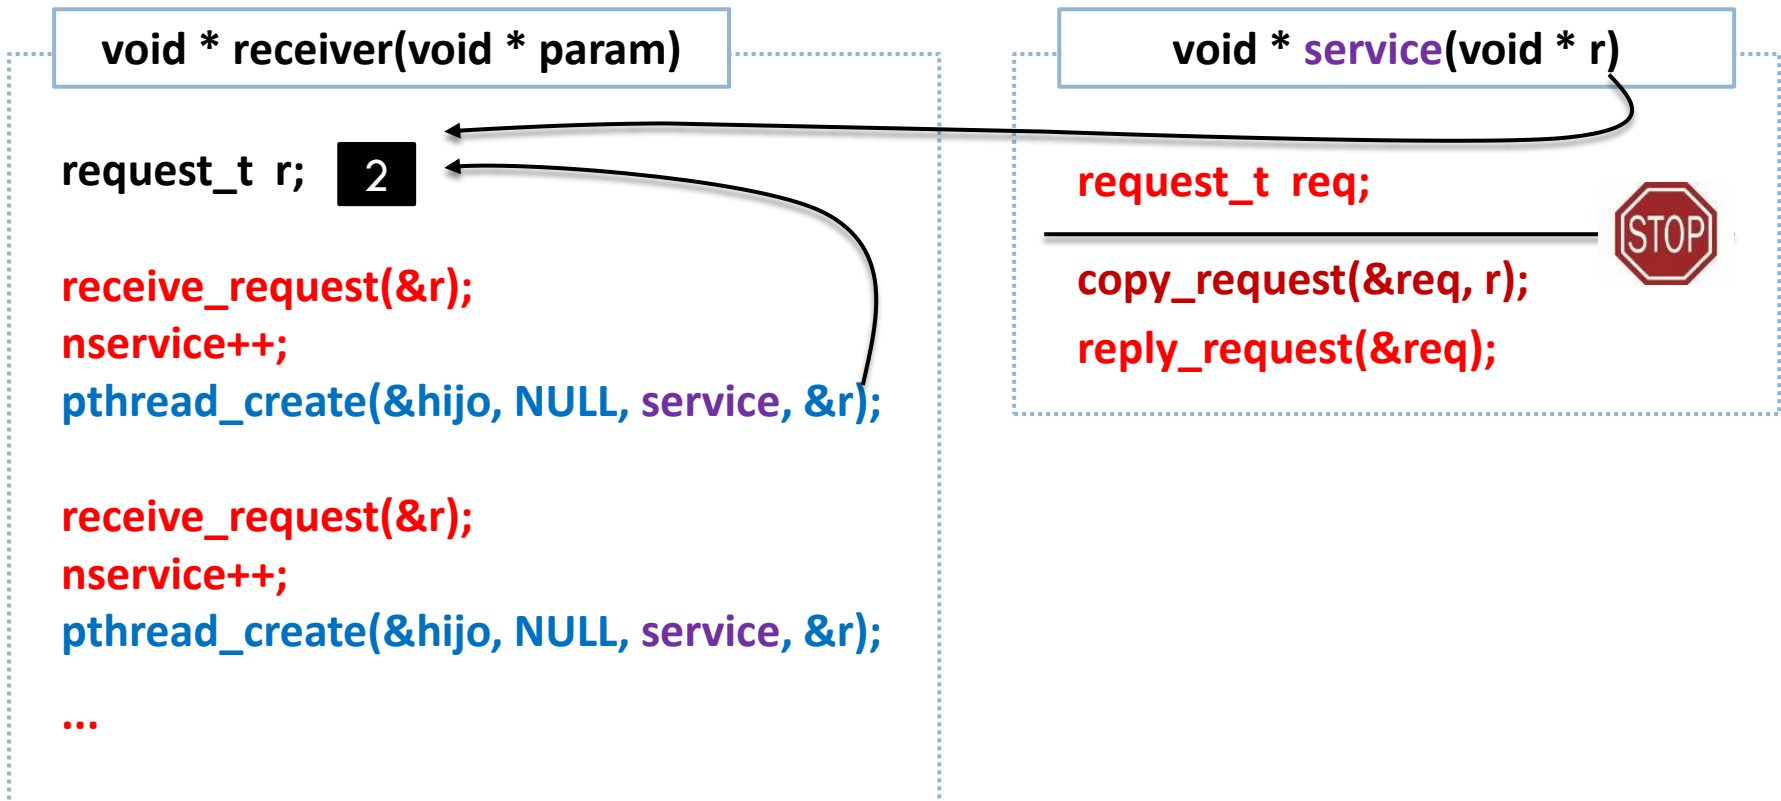
```
void * service(void * r)
```

```
request_t req;  
  
copy_request(&req, r);  
reply_request(&req);
```

Race condition

40

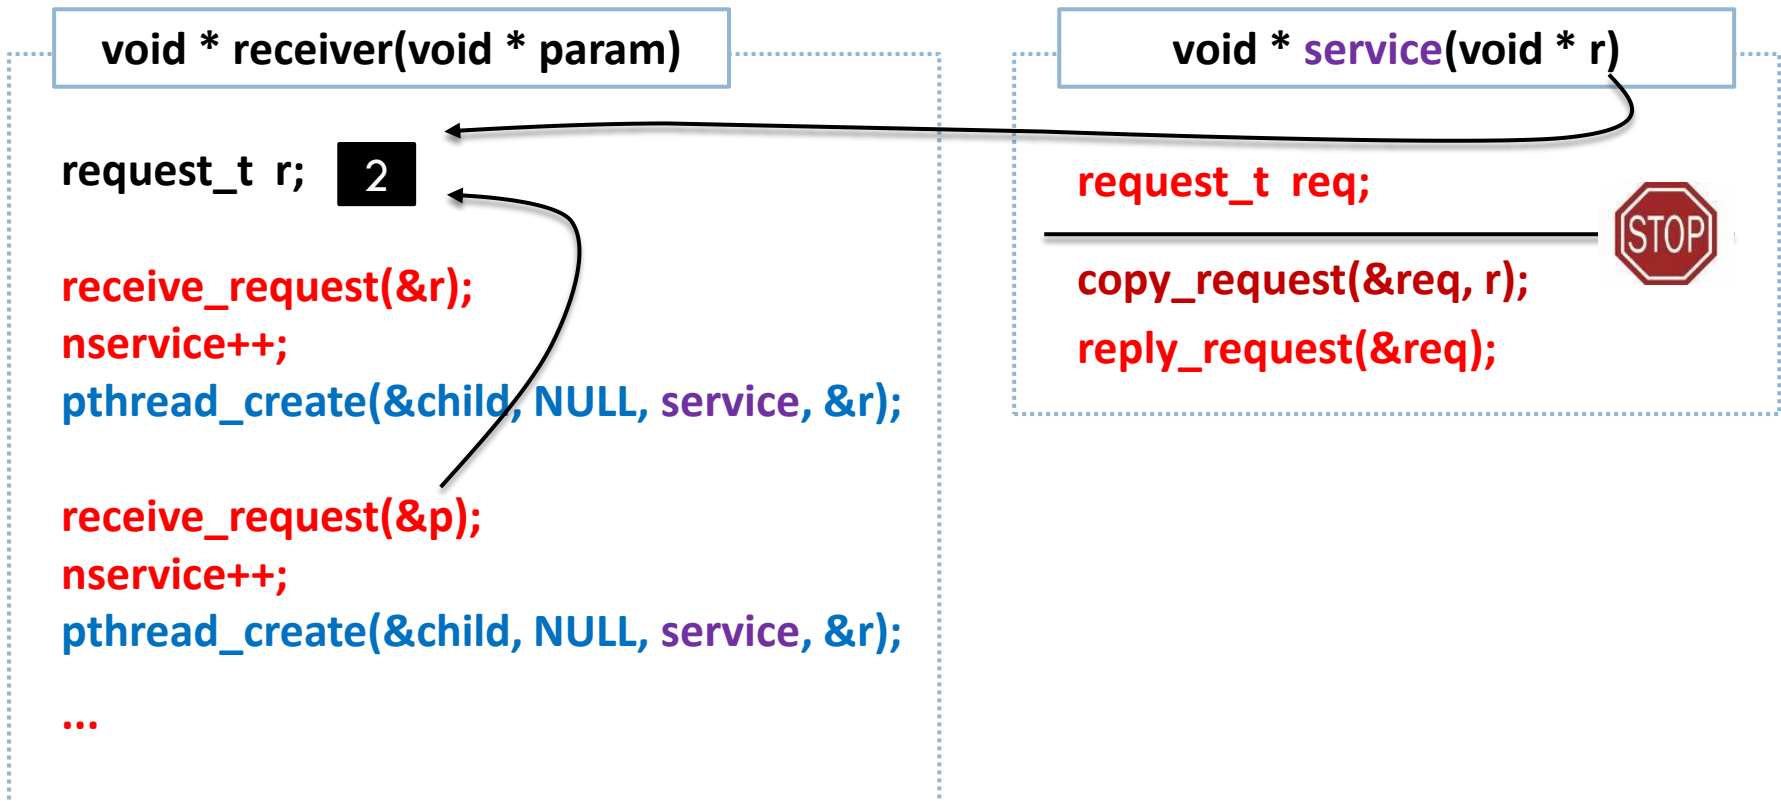
□ Can a race condition happen?



Race condition

41

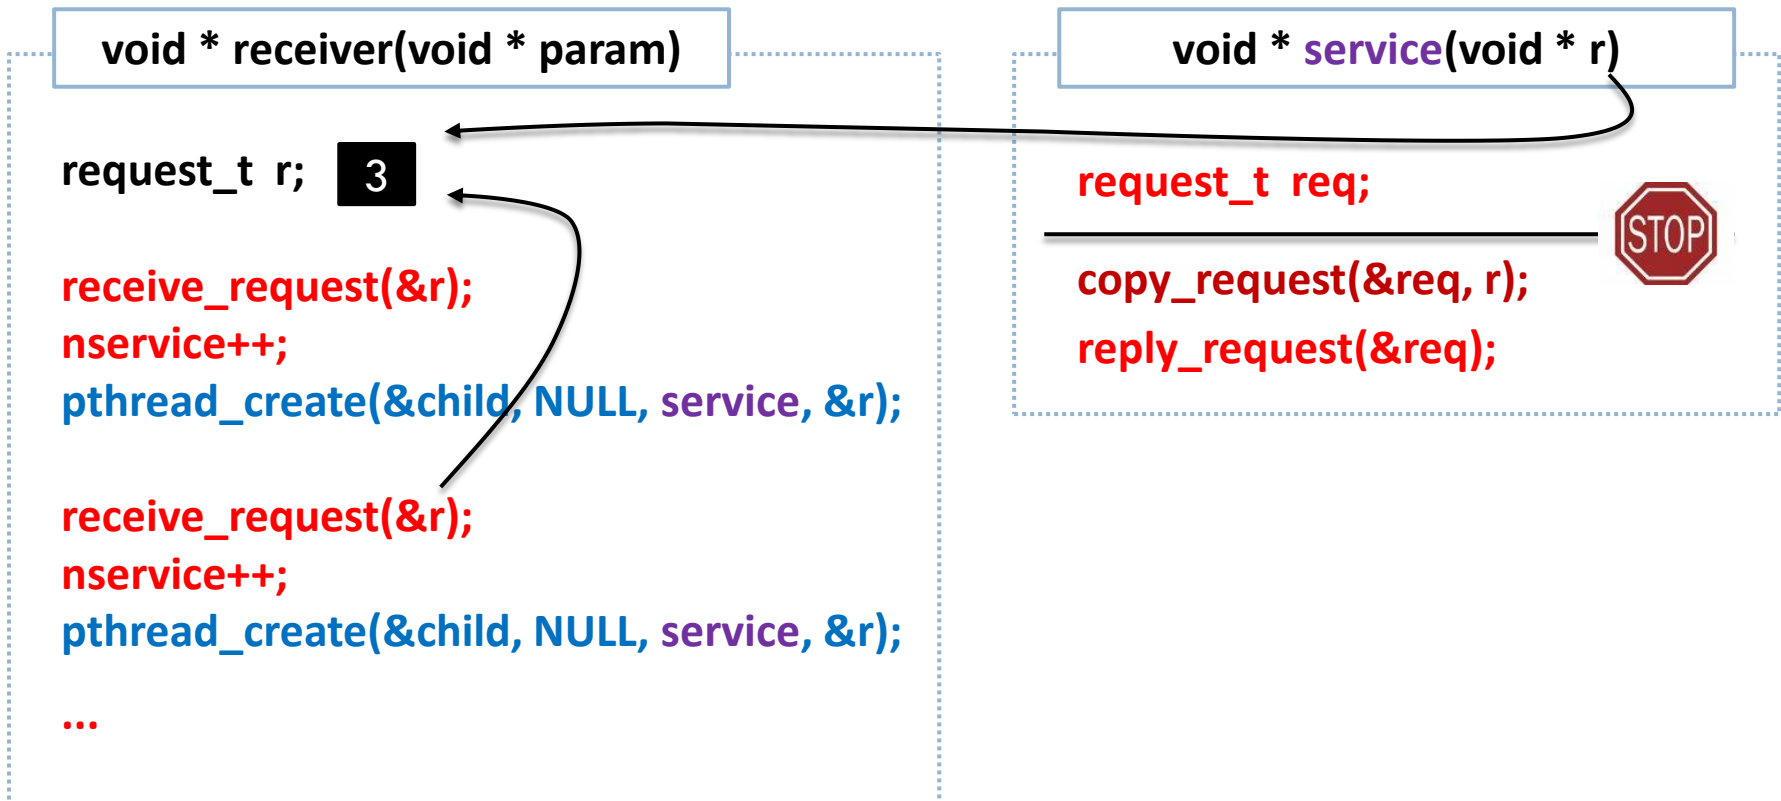
□ Can a race condition happen?



Race condition

42

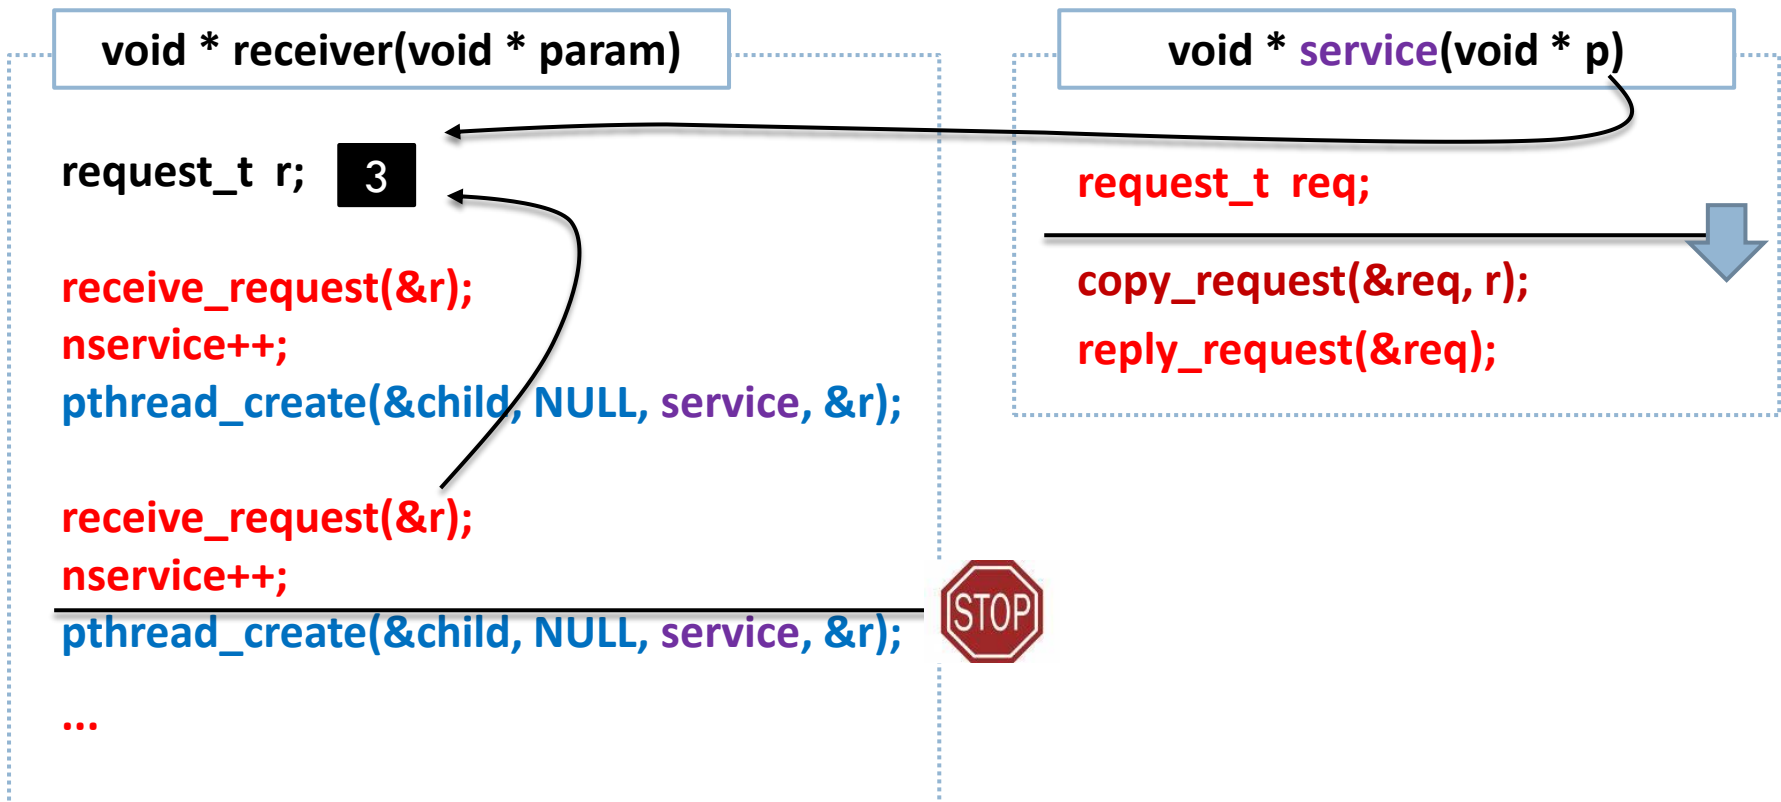
□ Can a race condition happen?



Race condition

43

□ Can a race condition happen?




Implementation

44

```
#include "request.h"
#include <stdio.h>
#include <time.h>
#include <pthread.h>
#include <semaphore.h>
```

```
sem_t snchildren;
sem_t sparam;
```

```
int main() {
    time_t t1, t2;
    double diff;
    pthread_t thr;
```



```
t1 = time(NULL);
sem_init(&snchildren, 0, 0);
sem_init(&sparam, 0, 0);
pthread_create(&thr, NULL,
               receiver, NULL);

pthread_join(thr, NULL);
sem_destroy(&snchildren);
sem_destroy(&sparam);
t2 = time(NULL);

diff = difftime(t2,t1);
printf("Time: %lf\n",diff);

return 0;
}
```



Implementation: receiver

45

```
void * receiver(void * param) {
    const int MAX_REQUESTS = 5;
    int i, nservice = 0;
    request_t r;
    pthread_t th_child;

    for (i=0;i<MAX_REQUESTS;i++) {
        receive_request(&p);
        nservice++;
        pthread_create(&th_child, NULL, service, &r);
        sem_wait(&sparam);
    }


    while (nservice>0) {
        fprintf(stderr, "Performing wait\n");
        sem_wait(&snchildren);
        nservice--;
        fprintf(stderr, "Exiting from wait\n");
    }

    pthread_exit(0);
    return NULL;
}
```

Implementation: service

46

```
void * service(void * r) {  
    request_t req;  
  
    copy_request(&req,(request_t*)r);  
    sem_post(&sparam);  
  
    fprintf(stderr, "Starting service\n");  
    reply_request(&req);  
    sem_post(&snchildren);  
  
    fprintf(stderr, "Finishing service\n");  
    pthread_exit(0);  
    return NULL;  
}
```



Evaluation

47

Normal	Processes	Thread per request	<i>Thread Pool</i>
54 seg.	17 seg.	20 seg.	

Problem

48

- Thread creation and termination has a cost which is lower than for processes,
 - ▣ But still a cost!

- No admission control:
 - ▣ What if many requests arrive or received requests do not finish?

Contents

49

- ❑ Request servers.
- ❑ Process based solution.
- ❑ Threads on demand solution.
- ❑ **Thread pool solution.**

Thread Pool

50

- A thread pool is a set of threads created at startup to run a service:
- Each time a request arrives it is placed in the queue of pending requests.
- All threads wait until some request is in the queue and they take it from the queue to process it.

Implementation: main (1 / 3)

51

```
#include "request.h"
#include <stdio.h>
#include <time.h>
#include <pthread.h>
#include <semaphore.h>

#define MAX_BUFFER 128
request_t buffer[MAX_BUFFER];
int n_elements;
int pos_service = 0;
pthread_mutex_t mutex;
pthread_cond_t non_full;
pthread_cond_t non_empty;

pthread_mutex_t mend;
int end=0;
```

Implementation: main (2/3)

52

```
int main()
{
    time_t t1, t2;
    double diff;
    pthread_t thr;
    const int MAX_SERVICE = 5; int i;
    pthread_t ths[MAX_SERVICE];

    t1 = time(NULL);

    pthread_mutex_init(&mutex, NULL);
    pthread_cond_init(&non_full, NULL);
    pthread_cond_init(&non_empty, NULL);
    pthread_mutex_init(&mend, NULL);

    pthread_create(&thr, NULL, receiver, NULL);
    for (i=0; i<MAX_SERVICE; i++) {
        pthread_create(&ths[i], NULL, service, NULL);
    }
}
```

Implementation: main (3/3)

53

```
pthread_join(thr, NULL);
for (i=0;i<MAX_SERVICE;i++) {
    pthread_join(thr[i],NULL);
}

pthread_mutex_destroy(&mutex);
pthread_cond_destroy(&non_full);
pthread_cond_destroy(&non_empty);
pthread_mutex_destroy(&mend);

t2 = time(NULL);

diff = difftime(t2,t1);
printf("Time: %lf\n",diff);

return 0;
}
```

Implementation: receiver (1 / 2)

54

```
void * receiver (void * param)
{
    const int MAX_REQUESTS = 5;
    request_t r;
    int i, pos=0;

    for (i=0;i<MAX_REQUESTS;i++)
    {
        receive_request(&r);
        pthread_mutex_lock(&mutex);
        while (n_elements == MAX_BUFFER)
            pthread_cond_wait(&non_full, &mutex);
        buffer[pos] = r;
        pos = (pos+1) % MAX_BUFFER;
        n_elements++;
        pthread_cond_signal(&non_empty);
        pthread_mutex_unlock(&mutex);
    }
}
```

Implementation: receiver (2/2)

55

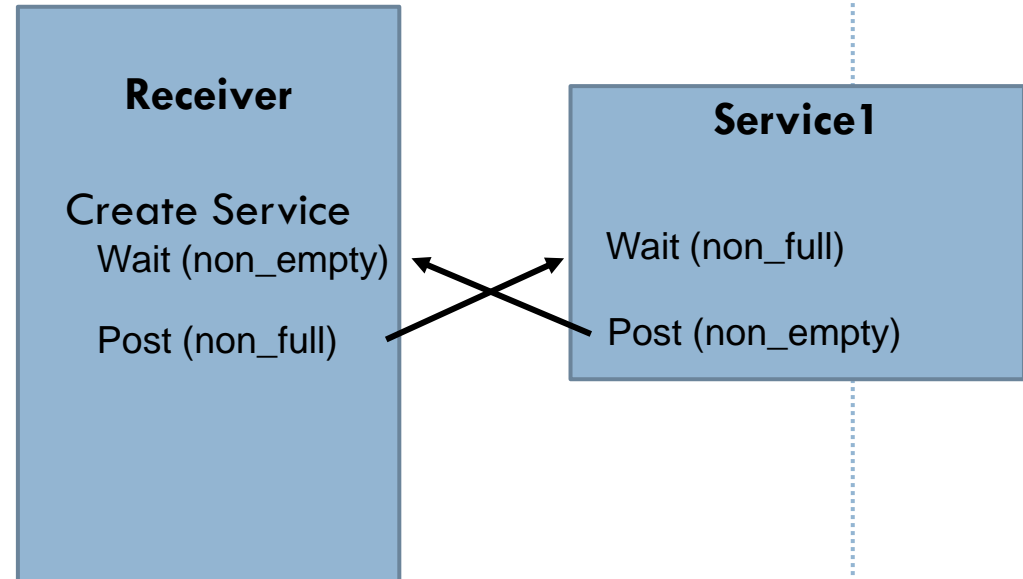
```
fprintf(stderr, "Finishing receiver\n");  
pthread_mutex_lock(&mend);  
end=1;  
pthread_mutex_unlock(&mend);  
  
pthread_mutex_lock(&mutex);  
pthread_cond_broadcast(&non_empty);  
pthread_mutex_unlock(&mutex);  
  
fprintf(stderr, "Finished receiver\n");  
pthread_exit(0);  
return NULL;  
  
} /* receiver*/
```

Implementation: service (1 / 2)

56

```
void * service (void * param)
{
    request_t r;

    for (;;) {
        pthread_mutex_lock(&mutex);
        while (n_elements == 0) {
            if (end==1) {
                fprintf(stderr, "Finishing service\n");
                pthread_mutex_unlock(&mutex);
                pthread_exit(0);
            }
            pthread_cond_wait(&non_empty, &mutex);
        } // while
    }
```



Implementation: service (2/2)

57

```
fprintf(stderr, "Serving position %d\n", pos_service);
r = buffer[pos_service];
pos_service = (pos_service + 1) % MAX_BUFFER;
n_elements --;
pthread_cond_signal(&non_full);
pthread_mutex_unlock(&mutex);
reply_request(&r);
}

pthread_exit(0);
return NULL;
}
```

Evaluation

58

Normal	Processes	Thread per request	<i>Thread Pool</i>
54 seg.	17 seg.	20 seg.	12 seg